
Distributed Learning: Risk Bound and Algorithm

Anonymous Author(s)

Affiliation

Address

email

Abstract

1

2 1 Preliminaries

3 We consider the supervised learning where a learning algorithm receives a sample of N labeled points

$$\mathcal{S} = \{z_i = (\mathbf{x}_i, y_i)\}_{i=1}^N \in (\mathcal{Z} = \mathcal{X} \times \mathcal{Y})^N,$$

4 where \mathcal{X} denotes the input space and \mathcal{Y} denotes the output space. We assume \mathcal{S} is drawn identically
5 and independently from a fixed, but unknown probability distribution \mathbb{P} on $\mathcal{Z} = \mathcal{X} \times \mathcal{Y}$. The goal is
6 to learn a good prediction model $f \in \mathcal{H} : \mathcal{X} \rightarrow \mathcal{Y}$, whose prediction accuracy at instance $z = (\mathbf{x}, y)$
7 is measured by a loss function $\ell(f, z)$.

8 In this paper, we focus on the supervised learning over the some Hilbert space \mathcal{H} :

$$\hat{f} = \arg \min_{f \in \mathcal{H}} \hat{R}(f) = \frac{1}{N} \sum_{i=1}^N \ell(f, z_i) + r(f) \quad (1)$$

9 where $\ell(f, z)$ is the loss function, and $r(f)$ is a regularizer.

10 The expect of \hat{R} is defined as

$$R(f) = \mathbb{E}_{z \sim \mathbb{P}}[\ell(f, z)] + r(f). \quad (2)$$

11 Let $f^* = \arg \min_{f \in \mathcal{H}} R(f)$, and $R_* = R(f^*)$.

12 In the distributed setting, we divide evenly amongst m processors or inference procedures. Let
13 $\mathcal{S}_i, i \in (1, 2, \dots, m)$, denote a subsampled dataset of size $n = \frac{N}{m}$. For each $i = 1, 2, \dots, m$, the
14 local estimate

$$\hat{f}_i = \arg \min_{f \in \mathcal{H}} \hat{R}_i(f) = \left\{ \frac{1}{n} \sum_{z_i \in \mathcal{S}_i} \ell(f, z_i) + r(f) \right\}$$

15 The average local estimates is denote as

$$\bar{f} = \frac{1}{m} \sum_{i=1}^m \hat{f}_i.$$

16 In the next, we will estimate the discrepancy of $R(\bar{f})$ and $R(f^*)$.

17 2 Faster Rates of Distributed Learning

18 2.1 Assumptions

19 In the following, we use $\|\cdot\|_{\mathcal{H}}$ to denote the norm induced by inner product of the Hilbert space \mathcal{H} .

Assumption 1. The function $\nu(f, z) = \ell(f, z) + r(f)$ is η -strongly convex and β -smooth with respect to the first variable f , that is $\forall f, f' \in \mathcal{H}, z \in \mathcal{Z}$,

$$\langle \nabla \nu(f, z), f - f' \rangle_{\mathcal{H}} + \frac{\eta}{2} \|f - f'\|_{\mathcal{H}} \leq \nu(f, z) - \nu(f', z), \quad (3)$$

$$\|\nabla \nu(f, z) - \nabla \nu(f', z)\|_{\mathcal{H}} \leq \beta \|f - f'\|_{\mathcal{H}}. \quad (4)$$

The above assumptions allow us to model many popular losses, such as square loss and logistic loss, and the regularizer, $r(f) = \lambda \|f\|_{\mathcal{H}}^2$.

Assumption 2. Let $f_* = \arg \min_{f \in \mathcal{H}} R(f)$. We assume that the gradient at f_* is upper bounded by M , that is

$$\|\nabla \ell(f_*, z)\|_{\mathcal{H}} \leq M, \forall z \in \mathcal{Z}.$$

The above assumption is a common assumption can be seen in [4, 2].

3 Faster Rates of Distributed Learning

Theorem 1. For any $0 < \delta < 1$, ϵ , the cover number of the Hilbert space \mathcal{H} is defined as $\mathcal{N}(\mathcal{H}, \epsilon)$. Under Assumptions 1 and 2, and if

$$m \leq \frac{N\eta}{4\beta \log \mathcal{N}(\mathcal{H}, \epsilon)}, \quad (5)$$

with probability at least $1 - \delta$, we have

$$\begin{aligned} R(\bar{f}) - R(f_*) &\leq \frac{16\beta \log(4m/\delta)}{n^2\eta} + \frac{128\beta R_* \log(4m/\delta)}{n\eta} + \frac{32\beta^2 \epsilon^2}{\eta} + \frac{64\beta L \log(\mathcal{N}(\mathcal{H}, \epsilon)) \epsilon}{n\eta} \\ &\quad - \frac{64\beta \log^2(\mathcal{N}(\mathcal{H}, \epsilon)) \epsilon^2}{n^2\eta} - \Delta(\bar{f}), \end{aligned} \quad (6)$$

where $R_* = R(f_*)$, $\Delta_{\bar{f}} = \frac{\eta}{4m^2} \sum_{i,j=1, i \neq j}^m \|\hat{f}_i - \hat{f}_j\|_{\mathcal{H}}^2$

By choosing ϵ small enough,

$$\frac{32\beta^2 \epsilon^2}{\eta} + \frac{64\beta L \log(\mathcal{N}(\mathcal{H}, \epsilon)) \epsilon}{n\eta} + \frac{64\beta \log^2(\mathcal{N}(\mathcal{H}, \epsilon)) \epsilon^2}{n^2\eta}$$

will becomes non-dominating. To be specific, we have the following corollary:

Corollary 1. By setting $\epsilon = \frac{1}{n}$ in Theorem 1, with high probability, we have

$$R(\bar{f}) - R(f_*) = \mathcal{O} \left(\frac{R_* \log(m)}{n} + \frac{\log(\mathcal{N}(\mathcal{H}, \frac{1}{n}))}{n^2} - \Delta(\bar{f}) \right).$$

3.1 Linear Hypothesis Space

If we consider the linear hypothesis space, that is

$$\mathcal{H} = \{f = \mathbf{w}^T \mathbf{x} | \mathbf{w} \in \mathbb{R}^d, \|\mathbf{w}\|_2 \leq B\}.$$

According to the [1], the cover number of linear hypothesis space can be bounded:

$$\log(\mathcal{N}(\mathcal{H}, 1/n)) \leq d \log(6Bn).$$

Thus, from Corollary 1, we have

$$R(\bar{f}) - R(f_*) = \mathcal{O} \left(\frac{R_* \log m}{n} + \frac{d \log n}{n^2} - \Delta(\bar{f}) \right)$$

When the minimal risk is small, i.e., $R_* = \mathcal{O}(\frac{d}{n})$, the rate is improved to

$$\mathcal{O} \left(\frac{d \log(mn)}{n^2} - \Delta(\bar{f}) \right) = \mathcal{O} \left(\frac{d \log N}{n^2} - \Delta(\bar{f}) \right).$$

Therefore, if $m \leq \sqrt{\frac{N}{d \log N}}$, we have

$$R(\bar{f}) - R(f_*) = \mathcal{O} \left(\frac{1}{N} - \Delta(\bar{f}) \right).$$

3.2 Reproducing Kernel Hilbert Space

The reproducing kernel Hilbert space \mathcal{H}_K associated with the kernel K is defined to be the closure of the linear span of the set of functions $\{K(\mathbf{x}, \cdot) : \mathbf{x} \in \mathcal{X}\}$ with the inner product satisfying

$$\langle K(\mathbf{x}, \cdot), f \rangle_{\mathcal{H}_K} = f(\mathbf{x}), \forall \mathbf{x} \in \mathcal{X}, f \in \mathcal{H}_K.$$

In this subsection, we consider hypothesis space as the reproducing kernel Hilbert space,

$$\mathcal{H} := \{f \in \mathcal{H}_K : \|f\|_{\mathcal{H}_K} \leq B\}.$$

From [5], if the Mercer kernel

$$K(\mathbf{x}, \mathbf{x}') = k(\mathbf{x} - \mathbf{x}'), k(\mathbf{x}) = \exp \left\{ -\frac{\|\mathbf{x}\|^2}{\sigma^2}, \mathbf{x}, \mathbf{x}' \in [0, 1]^d, \right\},$$

then for $0 \leq \epsilon \leq B/2$, there holds:

$$\log(\mathcal{N}(\mathcal{H}, 1/n)) = \mathcal{O}(\log^d(nB))$$

According to Corollary 1, we can obtain that

$$R(\bar{f}) - R(f_*) = \mathcal{O} \left(\frac{R_* \log m}{n} + \frac{\log^d n}{n^2} - \Delta(\bar{f}) \right).$$

When the minimal risk R_* is small, $R_* = \mathcal{O} \left(\frac{\log^{(d-1)} n}{n} \right)$, if $m \leq n$, we have

$$R(\bar{f}) - R(f_*) = \mathcal{O} \left(\frac{\log^d n}{n^2} - \Delta(\bar{f}) \right)$$

Therefore, if $m \leq \sqrt{\frac{N}{\log^2 n}}$, we have

$$R(f) - R(f_*) = \mathcal{O} \left(\frac{1}{N} - \Delta(\bar{f}) \right)$$

3.3 Comparison with Related Work

Under the smooth, strongly convex and other some common assumption, [4] shows that

$$\mathbb{E} [\|\bar{f} - f_*\|^2] = \mathcal{O} \left(\frac{1}{N} + \frac{\log d}{n^2} \right). \quad (7)$$

If $\nu(f, z)$ is L -Lipschitz continuous over f , that is

$$\forall f, f' \in \mathcal{H}, z \in \mathcal{Z}, |\nu(f, z) - \nu(f', z)| \leq L \|f - f'\|_{\mathcal{H}},$$

it is easy to verify that

$$\begin{aligned} R(f) - R(f_*) &\leq L \mathbb{E} [\|\bar{f} - f_*\|_{\mathcal{H}}] \leq L \sqrt{\mathbb{E} [\|\bar{f} - f_*\|_{\mathcal{H}}^2]} \\ &= \mathcal{O} \left(\frac{1}{\sqrt{N}} + \frac{\sqrt{\log d}}{n} \right). \end{aligned} \quad (8)$$

According to the subsections and , we know that if m is not very large, the order of this paper can reach $\mathcal{O} \left(\frac{1}{N} - \Delta(\bar{f}) \right)$, which is much sharper than the order of (8).

[3] consider the kernel ridge regression, under some assumptions over the feature map induced by the kernel function, and if m is not very large, they show that

$$\mathbb{E} [\|\bar{f} - f_*\|^2] = \mathcal{O} \left(\frac{1}{N} \right).$$

If $\nu(f, z)$ is L -Lipschitz continuous over f , same as the above analysis, it is easy to verify that

$$R(f) - R(f_*) = \mathcal{O} \left(\frac{1}{\sqrt{N}} \right),$$

which is much looser than of propose bound.

60 4 Discrepant Distributed Algorithms (DDA)

61 According to the above results, under some assumptions, we know that

$$R(f) - R(f_*) = \mathcal{O}\left(\frac{1}{N} - \Delta(\bar{f})\right),$$

62 where $\Delta_{\bar{f}} = \frac{\eta}{4m^2} \sum_{i,j=1, i \neq j}^m \|\hat{f}_i - \hat{f}_j\|_{\mathcal{H}}^2$. Thus, to obtain tight bound, the discrepancy of each
 63 local estimate $\hat{f}_i, i = 1, \dots, m$ should be large. Therefore, it is reasonable to derive the following
 64 optimization problem:

$$\hat{f}_i = \arg \min_{f \in \mathcal{H}} \frac{1}{n} \sum_{z_j \in \mathcal{S}_i} \ell(f, z_j) + r(f) - \gamma \|f - \bar{f}_{\setminus i}\|_{\mathcal{H}}, \quad (9)$$

65 where $\bar{f}_{\setminus i} = \frac{1}{m-1} \sum_{j=1, j \neq i}^m \hat{f}_j$.

66 4.1 Linear Hypothesis Space

67 When \mathcal{H} is a linear Hypothesis space, we consider the following problem:

$$\hat{\mathbf{w}}_i = \arg \min_{\mathbf{w} \in \mathbb{R}^d} \frac{1}{n} \sum_{z_i \in \mathcal{S}_i} (\mathbf{w}^T \mathbf{x}_i - y_i)^2 + \lambda \|\mathbf{w}\|_2^2 - \gamma \|\mathbf{w} - \bar{\mathbf{w}}_{\setminus i}\|_2^2,$$

68 where $\bar{\mathbf{w}}_{\setminus i} = \frac{1}{m-1} \sum_{j=1, j \neq i} \hat{\mathbf{w}}_j$. If given $\bar{\mathbf{w}}_{\setminus i} = \frac{1}{m-1} \sum_{j=1, j \neq i} \hat{\mathbf{w}}_j$, it is easy to verify that $\hat{\mathbf{w}}_i$ can
 69 be written as

$$\hat{\mathbf{w}}_i = \left(\frac{1}{n} \mathbf{X}_{\mathcal{S}_i} \mathbf{X}_{\mathcal{S}_i}^T + \lambda \mathbf{I}_d - \gamma \mathbf{I}_d \right)^{-1} \left(\frac{1}{n} \mathbf{X}_{\mathcal{S}_i} \mathbf{y}_{\mathcal{S}_i} - \gamma \bar{\mathbf{w}}_{\setminus i} \right),$$

70 where $\mathbf{X}_{\mathcal{S}_i} = (\mathbf{x}_{t_1}, \mathbf{x}_{t_2}, \dots, \mathbf{x}_{t_n})$, $\mathbf{y}_{\mathcal{S}_i} = (y_{t_1}, y_{t_2}, \dots, y_{t_n})^T$, $z_{t_i} \in \mathcal{S}_i, i = 1, \dots, n$.

71 Let $\mathbf{A}_i = \frac{1}{n} \mathbf{X}_{\mathcal{S}_i} \mathbf{X}_{\mathcal{S}_i}^T + \lambda \mathbf{I}_d - \gamma \mathbf{I}_d$, $\mathbf{b}_i = \frac{1}{n} \mathbf{X}_{\mathcal{S}_i}^T \mathbf{y}_{\mathcal{S}_i}$.

72 Let $\mathbf{d}_i = \mathbf{A}_i^{-1} \bar{\mathbf{w}}_{\setminus i}$, $\hat{\mathbf{w}}_i = \mathbf{A}_i^{-1} \mathbf{b}_i$, we have

$$\bar{\mathbf{w}}_{\setminus i}^T \hat{\mathbf{w}}_i = \bar{\mathbf{w}}_{\setminus i}^T \mathbf{A}_i^{-1} \mathbf{b}_i = (\mathbf{A}_i^{-1} \bar{\mathbf{w}}_{\setminus i})^T \mathbf{b}_i = \mathbf{d}_i^T \mathbf{b}_i,$$

73 thus $\mathbf{d}_i = \frac{\bar{\mathbf{w}}_{\setminus i}^T \hat{\mathbf{w}}_i}{\mathbf{b}_i}$

74 The DDA Algorithm is given as follows:

75 **Input** : $\lambda, \gamma, \mathbf{X}, m, \zeta > 0$.

76 **For** $t = 0, 1, \dots, T$

77 Each branch node i :

78 **If** $t = 0$

79 $\hat{\mathbf{w}}_i^0 = \mathbf{A}_i^{-1} \mathbf{b}_i$;

80 **else**

81 $\mathbf{d}_i^t = \frac{(\bar{\mathbf{w}}_{\setminus i}^0)^T \hat{\mathbf{w}}_i^0}{\mathbf{b}_i}$

82 $\hat{\mathbf{w}}_i^t = \hat{\mathbf{w}}_i^0 - \gamma \mathbf{d}_i^t$;

83 push $\hat{\mathbf{w}}_i^t$ to center node;

84 Center node:

85 $\bar{\mathbf{w}}^t = \frac{1}{m} \sum_{i=1}^m \hat{\mathbf{w}}_i^t$

86 **If** $\|\bar{\mathbf{w}}^t - \bar{\mathbf{w}}^{t-1}\| \leq \zeta$, **End**

87 **else** push $\bar{\mathbf{w}}_{\setminus i}^t = \frac{m \bar{\mathbf{w}}^t - \hat{\mathbf{w}}_i^t}{m-1}$ to each branch node i

88 **End**

89 **Output** : $\bar{\mathbf{w}} = \frac{1}{m} \sum_{i=1}^m \hat{\mathbf{w}}_i^T$

90 4.2 Reproducing Kernel Hilbert Space

91 When \mathcal{H} is a reproducing kernel Hilbert space, we consider the following problem:

$$\hat{f}_i = \arg \min_{f \in \mathcal{H}} \frac{1}{n} \sum_{z_i \in \mathcal{S}_i} (f(\mathbf{x}_i) - y_i)^2 + \lambda \|f\|_{\mathcal{H}}^2 + \gamma \|f - \bar{f}_i\|_{\mathcal{H}}^2, \quad (10)$$

92 where $f(\mathbf{x}) = \sum_{j=1}^n c_j K(\mathbf{x}_j, \mathbf{x})$, which can be written as

$$\hat{\mathbf{c}}_i = \arg \min_{\mathbf{c} \in \mathbb{R}^n} \frac{1}{n} \|\mathbf{K}_{\mathcal{S}_i} \mathbf{c} - \mathbf{y}_{\mathcal{S}_i}\|_2^2 + \lambda \mathbf{c}^T \mathbf{K}_{\mathcal{S}_i} \mathbf{c} - \gamma (\mathbf{c} - \bar{\mathbf{c}}_i)^T \mathbf{K}_{\mathcal{S}_i} (\mathbf{c} - \bar{\mathbf{c}}_i). \quad (11)$$

93 If given $\bar{\mathbf{c}}_i = \frac{1}{m-1} \sum_{j=1, j \neq i} \hat{\mathbf{c}}_j$, it is easy to verify that $\hat{\mathbf{c}}_i$ can be written as

$$\hat{\mathbf{c}}_i = (\mathbf{K}_{\mathcal{S}_i} + \lambda \mathbf{I}_n - \gamma \mathbf{I}_n)^{-1} (\mathbf{y}_{\mathcal{S}_i} - \gamma \bar{\mathbf{c}}_i)$$

94 Let $\mathbf{A}_i = \mathbf{K}_{\mathcal{S}_i} + \lambda \mathbf{I}_n - \gamma \mathbf{I}_n$, $\mathbf{b}_i = \mathbf{y}_{\mathcal{S}_i}$.

95 **Input** : $\lambda, \gamma, \mathbf{X}, m, \zeta > 0$.

96 **For** $t = 0, 1, \dots, T$

97 Each branch node i :

98 **If** $t = 0$

99 $\hat{\mathbf{c}}_i^0 = \mathbf{A}_i^{-1} \mathbf{b}_i$;

100 **else**

101 $\mathbf{d}_i^t = \frac{(\bar{\mathbf{c}}_i^0)^T \hat{\mathbf{c}}_i^0}{\mathbf{b}_i}$

102 $\hat{\mathbf{c}}_i^t = \hat{\mathbf{c}}_i^0 - \gamma \mathbf{d}_i^t$;

103 push $\hat{\mathbf{w}}_i^t$ to center node;

104 Center node:

105 $\bar{\mathbf{c}}^t = \frac{1}{m} \sum_{i=1}^m \hat{\mathbf{c}}_i^t$

106 **If** $\|\bar{\mathbf{c}}^t - \bar{\mathbf{c}}^{t-1}\| \leq \zeta$, **End**

107 **else** push $\bar{\mathbf{c}}_i^t = \frac{m\bar{\mathbf{c}}^t - \hat{\mathbf{c}}_i^t}{m-1}$ to each branch node i

108 **End**

109 **Output** : $\bar{\mathbf{c}} = \frac{1}{m} \sum_{i=1}^m \hat{\mathbf{c}}_i^T$

110 5 Analysis

111 5.1 The Key Idea

112 Since $R(f)$ is η -strongly convex function, we have

$$R(\bar{f}) \leq \frac{1}{m} \sum_{i=1}^m R(\hat{f}_i) - \frac{\eta}{2m^2} \sum_{i,j=1}^m \|\hat{f}_i - \hat{f}_j\|^2.$$

113 Therefore, we have

$$R(\bar{f}) - R(f_*) \leq \frac{1}{m} \sum_{i=1}^m [R(\hat{f}_i) - R(f_*)] - \frac{\eta}{4m^2} \sum_{i,j=1, i \neq j}^m \|\hat{f}_i - \hat{f}_j\|^2. \quad (12)$$

114 In the next, we will estimate $R(\hat{f}_i) - R(f_*)$.

115 Our theoretical analysis is built upon the following inequality:

$$\begin{aligned} & R(\hat{f}_i) - R(f_*) + \frac{\eta}{2} \|\hat{f}_i - f_*\|^2 \leq \langle \nabla R(\hat{f}_i), \hat{f}_i - f_* \rangle \\ &= \langle \nabla R(\hat{f}_i) - \nabla R(f_*) - [\nabla \hat{R}(\hat{f}_i) - \nabla \hat{R}(f_*)], \hat{f}_i - f_* \rangle + \langle \nabla \hat{R}(\hat{f}_i) - \nabla \hat{R}(f_*) + \nabla R(f_*), \hat{f}_i - f_* \rangle \\ &= \langle \nabla R(\hat{f}_i) - \nabla R(f_*) - [\nabla \hat{R}(\hat{f}_i) - \nabla \hat{R}(f_*)], \hat{f}_i - f_* \rangle + \langle \nabla R(f_*) - \nabla \hat{R}(f_*), \hat{f}_i - f_* \rangle \\ &\leq \left(\underbrace{\left\| \nabla R(\hat{f}_i) - \nabla R(f_*) - [\nabla \hat{R}(\hat{f}_i) - \nabla \hat{R}(f_*)] \right\|}_{:=A_1} + \underbrace{\left\| \nabla R(f_*) - \nabla \hat{R}(f_*) \right\|}_{:=A_2} \right) \|\hat{f}_i - f_*\| \end{aligned} \quad (13)$$

116 where η is the strong convexity modulus of $R(\cdot)$ if exists otherwise it is zero.

117 **Lemma 1.** Let \mathcal{H} be a Hilbert space and let ξ be a random variable with values in \mathcal{H} . Assume
 118 $\|\xi\| \leq M \leq \infty$ almost surely. Denote $\sigma^2(\xi) = \mathbb{E}[\|\xi\|^2]$. Let $\{\xi_i\}_{i=1}^n$ be n independent drawers of ξ .
 119 For any $0 \leq \delta \leq 1$, with confidence $1 - \delta$,

$$\left\| \frac{1}{n} \sum_{j=1}^n [\xi_j - \mathbb{E}[\xi_j]] \right\| \leq \frac{2M \log(2/\delta)}{n} + \sqrt{\frac{2\sigma^2(\xi) \log(2/\delta)}{n}}.$$

120 **Lemma 2.** Under Assumptions ??, ?? and ??, with probability at least $1 - \delta$, for any $f \in \mathcal{N}(\mathcal{H}, \epsilon)$,
 121 we have

$$\begin{aligned} & \left\| \nabla R(f) - \nabla R(f_*) - [\nabla \hat{R}(f) - \nabla \hat{R}(f_*)] \right\| \\ & \leq \frac{\beta \log(\mathcal{N}(\mathcal{H}, \epsilon)) \|f - f_*\|}{n} + \sqrt{\frac{\beta \log(\mathcal{N}(\mathcal{H}, \epsilon)) (R(f) - R(f_*))}{n}}. \end{aligned} \quad (14)$$

122 *Proof.* Let $\nu(f) = \ell(f, \cdot) + r(f)$. Note that $\nu(f)$ is β -smooth, so we have

$$\|\nabla \nu(f) - \nabla \nu(f_*)\| \leq \beta \|f - f_*\| \quad (15)$$

123 Because ν is β -smooth and convex, by (2.1.7) of?, we have

$$\|\nabla \nu(f) - \nabla \nu(f_*)\|^2 \leq \beta (\nu(f) - \nu(f_*) - \langle \nabla \nu(f_*), f - f_* \rangle).$$

124 Taking expectation over both sides, we have

$$\begin{aligned} & \mathbb{E}[\|\nabla \nu(f) - \nabla \nu(f_*)\|^2] \\ & \leq \beta \left(R(\hat{f}_i) - R(f_*) - \langle \nabla R(f_*), f - f_* \rangle \right) \\ & \leq \beta \left(R(\hat{f}_i) - R(f_*) \right) \end{aligned}$$

125 where the last inequality follows from the optimality condition of f_* , i.e.,

$$\langle \nabla R(f_*), f - f_* \rangle \geq 0, \forall f \in \mathcal{H}.$$

126

□

127 Following Lemma 1, with probability at least $1 - \delta$, we have

$$\begin{aligned} & \left\| \nabla R(f) - \nabla R(f_*) - [\nabla \hat{R}(f) - \nabla \hat{R}(f_*)] \right\| \\ & = \left\| \nabla R(f) - \nabla R(f_*) - \frac{1}{n} \sum_{z_i \in \mathcal{S}_i} [\nabla \nu(f) - \nabla \nu(f_*)] \right\| \\ & \leq \frac{2\beta \|f - f_*\| \log(2/\delta)}{n} + \sqrt{\frac{2\beta (R(f) - R(f_*)) \log(2/\delta)}{n}} \end{aligned}$$

128 We obtain Lemma 2 by taking the union bound over all $f \in \mathcal{N}(\mathcal{H}, \epsilon)$.

129 **Lemma 3.** with probability at least $1 - \delta$, we have

$$\left\| \nabla R(f_*) - \nabla \hat{R}(f_*) \right\| \leq \frac{2M \log(2/\delta)}{n} + \sqrt{\frac{8\beta R_* \log(2/\delta)}{n}}. \quad (16)$$

130 *Proof.* Let $\nu(f, z_i) = \ell(f, z_i) + r(f)$ Since $\nu(\cdot, z_i)$ is β -smooth and nonnegative, from Lemma 4 of
 131 Srebro et al. (2010), we have

$$\|\nabla \nu(f_*, z_i)\|^2 \leq 4\beta \nu(f_*, z_i)$$

132 and thus

$$\mathbb{E}_{z \sim \mathbb{P}} [\|\nabla \nu(f_*, z)\|^2] \leq 4\beta \mathbb{E}_{z \sim \mathbb{P}} [\nu(f_*, z)] = 4\beta R(f_*).$$

133 From the **Assumption**, we have $\nabla\|\nu(f_*, z)\| \leq M, \forall z \in \mathcal{Z}$. Then, according to Lemma 1, with
 134 probability at least $1 - \delta$, we have

$$\begin{aligned} \left\| \nabla R(f_*) - \nabla \hat{R}(f_*) \right\| &= \left\| \nabla R(f_*) - \frac{1}{n} \sum_{z_j \in \mathcal{S}_i} \nabla \nu(f_*, z_j) \right\| \\ &\leq \frac{2\beta \log(2/\delta)}{n} + \sqrt{\frac{8\beta R_* \log(2/\delta)}{n}}. \end{aligned}$$

135

□

136 **Theorem 2.** At least $1 - 2\delta$, if

$$n \geq \frac{4\beta \log(\mathcal{N}(\mathcal{H}, \epsilon))}{\eta}, \quad (17)$$

137 we have

$$\begin{aligned} R(\bar{f}) - R(f_*) &\leq \frac{16\beta \log(2m/\delta)}{n^2 \eta} + \frac{128\beta R_* \log(2m/\delta)}{n \eta} + \frac{32\beta^2 \epsilon^2}{\eta} \\ &\quad + \frac{64\beta L \log(\mathcal{N}(\mathcal{H}, \epsilon)) \epsilon}{n \eta} + \frac{64\beta \log^2(\mathcal{N}(\mathcal{H}, \epsilon)) \epsilon^2}{n^2 \eta} \\ &\quad - \frac{\eta}{4m^2} \sum_{i,j=1, i \neq j}^m \|\hat{f}_i - \hat{f}_j\|^2. \end{aligned} \quad (18)$$

Proof. From the property of ϵ -net, we know that there exists a point $\tilde{f} \in \mathcal{N}(\mathcal{H}, \epsilon)$ such that

$$\|\hat{f}_i - \tilde{f}\| \leq \epsilon.$$

138 According to **Assumption ??**, we have

$$\begin{aligned} &\left\| \nabla R(\hat{f}_i) - \nabla R(f_*) - [\nabla \hat{R}(\hat{f}_i) - \nabla \hat{R}(f_*)] \right\| \\ &\leq \left\| \nabla R(\tilde{f}) - \nabla R(f_*) - [\nabla \hat{R}(\tilde{f}) - \nabla \hat{R}(f_*)] \right\| + 2\beta\epsilon \\ &\stackrel{(14)}{\leq} \frac{\beta \log(\mathcal{N}(\mathcal{H}, \epsilon)) \|\tilde{f} - f_*\|}{n} + \sqrt{\frac{\beta \log(\mathcal{N}(\mathcal{H}, \epsilon)) (R(\tilde{f}) - R(f_*))}{n}} + 2\beta\epsilon \\ &\leq \frac{\beta \log(\mathcal{N}(\mathcal{H}, \epsilon)) \|\hat{f}_i - f_*\|}{n} + \frac{\beta \log(\mathcal{N}(\mathcal{H}, \epsilon)) \epsilon}{n} + 2\beta\epsilon \\ &\quad + \sqrt{\frac{\beta \log(\mathcal{N}(\mathcal{H}, \epsilon)) (R(\hat{f}_i) - R(f_*))}{n}} + \sqrt{\frac{\beta \log(\mathcal{N}(\mathcal{H}, \epsilon)) (|R(\hat{f}_i) - R(\tilde{f})|)}{n}} \\ &\stackrel{(?)}{\leq} \frac{\beta \log(\mathcal{N}(\mathcal{H}, \epsilon)) \|\hat{f}_i - f_*\|}{n} + \frac{\beta \log(\mathcal{N}(\mathcal{H}, \epsilon)) \epsilon}{n} + 2\beta\epsilon \\ &\quad + \sqrt{\frac{\beta \log(\mathcal{N}(\mathcal{H}, \epsilon)) (R(\hat{f}_i) - R(f_*))}{n}} + \sqrt{\frac{\beta L \log(\mathcal{N}(\mathcal{H}, \epsilon)) \epsilon}{n}} \end{aligned} \quad (19)$$

139 Substituting (19) and (16) into (13), with probability at least $1 - 2\delta$, we have

$$\begin{aligned} &R(\hat{f}_i) - R(f_*) + \frac{\eta}{2} \|\hat{f}_i - f_*\|^2 \\ &\leq \frac{\beta \log(\mathcal{N}(\mathcal{H}, \epsilon)) \|\hat{f}_i - f_*\|^2}{n} + \frac{\beta \log(\mathcal{N}(\mathcal{H}, \epsilon)) \epsilon \|\hat{f}_i - f_*\|}{n} + 2\beta\epsilon \|\hat{f}_i - f_*\| \\ &\quad + \|\hat{f}_i - f_*\| \sqrt{\frac{\beta \log(\mathcal{N}(\mathcal{H}, \epsilon)) (R(\hat{f}_i) - R(f_*))}{n}} + \|\hat{f}_i - f_*\| \sqrt{\frac{\beta L \log(\mathcal{N}(\mathcal{H}, \epsilon)) \epsilon}{n}} \\ &\quad + \frac{2\beta \log(2/\delta) \|\hat{f}_i - f_*\|}{n} + \|\hat{f}_i - f_*\| \sqrt{\frac{8\beta R_* \log(2/\delta)}{n}}. \end{aligned} \quad (20)$$

140 Note that

$$\sqrt{ab} \leq \frac{a}{2\eta} + \frac{b\eta}{2}, \forall a, b, \eta \geq 0.$$

141 Therefore, we can obtain that

$$\begin{aligned} \|\hat{f}_i - f_*\| \sqrt{\frac{\beta \log(\mathcal{N}(\mathcal{H}, \epsilon)) (R(\hat{f}_i) - R(f_*))}{n}} &\leq \frac{2\beta \log(\mathcal{N}(\mathcal{H}, \epsilon)) (R(\hat{f}_i) - R(f_*))}{n\eta} + \frac{\eta}{8} \|\hat{f}_i - f_*\|^2, \\ \frac{2\beta \log(2/\delta) \|\hat{f}_i - f_*\|}{n} &\leq \frac{8\beta \log(2/\delta)}{n^2\eta} + \frac{\eta}{16} \|\hat{f}_i - f_*\|^2, \\ \|\hat{f}_i - f_*\| \sqrt{\frac{8\beta R_* \log(2/\delta)}{n}} &\leq \frac{64\beta R_* \log(2/\delta)}{n\eta} + \frac{\eta}{32} \|\hat{f}_i - f_*\|^2, \\ 2\beta\epsilon \|\hat{f}_i - f_*\| &\leq \frac{32\beta^2\epsilon^2}{\eta} + \frac{\eta}{64} \|\hat{f}_i - f_*\|^2, \\ \|\hat{f}_i - f_*\| \sqrt{\frac{\beta L \log(\mathcal{N}(\mathcal{H}, \epsilon)) \epsilon}{n}} &\leq \frac{32\beta L \log(\mathcal{N}(\mathcal{H}, \epsilon)) \epsilon}{n\eta} + \frac{\eta}{128} \|\hat{f}_i - f_*\|^2 \\ \frac{\beta \log(\mathcal{N}(\mathcal{H}, \epsilon)) \epsilon \|\hat{f}_i - f_*\|}{n} &\leq \frac{32\beta \log(\mathcal{N}(\mathcal{H}, \epsilon))^2 \epsilon^2}{n^2\eta} + \frac{\eta}{128} \|\hat{f}_i - f_*\|^2. \end{aligned}$$

142 Substituting the above inequation into (20), we can obtain that

$$\begin{aligned} &R(\hat{f}_i) - R(f_*) + \frac{\eta}{4} \|\hat{f}_i - f_*\|^2 \\ &\leq \frac{\beta \log(\mathcal{N}(\mathcal{H}, \epsilon)) \|\hat{f}_i - f_*\|^2}{n} + \frac{2\beta \log(\mathcal{N}(\mathcal{H}, \epsilon)) (R(\hat{f}_i) - R(f_*))}{n\eta} + \frac{8\beta \log(2/\delta)}{n^2\eta} \\ &\quad + \frac{64\beta R_* \log(2/\delta)}{n\eta} + \frac{32\beta^2\epsilon^2}{\eta} + \frac{32\beta L \log(\mathcal{N}(\mathcal{H}, \epsilon)) \epsilon}{n\eta} + \frac{32\beta \log^2(\mathcal{N}(\mathcal{H}, \epsilon)) \epsilon^2}{n^2\eta} \\ &\stackrel{(17)}{\leq} \frac{\eta}{4} \|\hat{f}_i - f_*\|^2 + \frac{1}{2} (R(\hat{f}_i) - R(f_*)) + \frac{8\beta \log(2/\delta)}{n^2\eta} \\ &\quad + \frac{64\beta R_* \log(2/\delta)}{n\eta} + \frac{32\beta^2\epsilon^2}{\eta} + \frac{32\beta L \log(\mathcal{N}(\mathcal{H}, \epsilon)) \epsilon}{n\eta} + \frac{32\beta \log^2(\mathcal{N}(\mathcal{H}, \epsilon)) \epsilon^2}{n^2\eta}. \end{aligned}$$

143 Thus, with $1 - 2\delta$, we have

$$\begin{aligned} R(\hat{f}_i) - R(f_*) &\leq \frac{16\beta \log(2/\delta)}{n^2\eta} + \frac{128\beta R_* \log(2/\delta)}{n\eta} + \frac{32\beta^2\epsilon^2}{\eta} \\ &\quad + \frac{64\beta L \log(\mathcal{N}(\mathcal{H}, \epsilon)) \epsilon}{n\eta} + \frac{64\beta \log^2(\mathcal{N}(\mathcal{H}, \epsilon)) \epsilon^2}{n^2\eta}. \end{aligned} \tag{21}$$

144 Combining (12) and (21), with $1 - 2\delta$, we have

$$\begin{aligned} R(\bar{f}) - R(f_*) &\leq \frac{16\beta \log(2m/\delta)}{n^2\eta} + \frac{128\beta R_* \log(2m/\delta)}{n\eta} + \frac{32\beta^2\epsilon^2}{\eta} + \frac{64\beta L \log(\mathcal{N}(\mathcal{H}, \epsilon)) \epsilon}{n\eta} + \\ &\quad + \frac{64\beta \log^2(\mathcal{N}(\mathcal{H}, \epsilon)) \epsilon^2}{n^2\eta} - \frac{\eta}{4m^2} \sum_{i,j=1, i \neq j}^m \|\hat{f}_i - \hat{f}_j\|^2. \end{aligned}$$

145

□

146 References

- 147 [1] G. Pisier. *The volume of convex bodies and Banach space geometry*, volume 94. Cambridge
148 University Press, 1999.
- 149 [2] L. Zhang, T. Yang, and R. Jin. Empirical risk minimization for stochastic convex optimization:
150 $O(1/n)$ -and $O(1/n^2)$ -type of risk bounds. In *Proceedings of the Conference on Learning Theory*
151 *(COLT 2017)*, pages 1954–1979, 2017.

- 152 [3] Y. Zhang, J. Duchi, and M. Wainwright. Divide and conquer kernel ridge regression. In
153 *Proceedings of Conference on Learning Theory (COLT 2013)*, pages 592–617, 2013.
- 154 [4] Y. Zhang, M. J. Wainwright, and J. C. Duchi. Communication-efficient algorithms for statistical
155 optimization. In *Advances in Neural Information Processing Systems*, pages 1502–1510, 2012.
- 156 [5] D.-X. Zhou. The covering number in learning theory. *Journal of Complexity*, 18(3):739–767,
157 2002.