
Max-Discrepancy Distributed Learning: Fast Risk Bounds and Algorithms

Anonymous Author(s)

Affiliation

Address

email

Abstract

We study the risk performance of distributed learning for the regularization empirical risk minimization with fast convergence rate, substantially improving the existing divide-and-conquer based distributed learning error analysis. An interesting theoretical finding is that the larger the discrepancy of each local estimate is, the tighter the risk bound is. This theoretical analysis motivates us to devise an effective max-discrepancy distributed learning algorithm (MDD). Experimental results show that our proposed method can outperform the existing divide-and-conquer methods but with little additional time cost. Theoretical analysis and empirical results demonstrate that our MDD is sound and effective.

1 Introduction

In the era of big data, the rapid expansion of computing capacities in automatic data generation and acquisition brings data of unprecedented size and complexity, and raises a series of scientific challenges such as storage bottleneck and algorithmic scalability [18, 15, 7]. Distributed learning based on a divide-and-conquer approach has triggered enormous recent research activities in various areas such as optimization [16] data mining [13] and machine learning [3]. This learning strategy breaks up a big problem into manageable pieces, operates learning algorithms on each piece on individual machines or processors, and then puts the individual solutions together to get a final global output. In this way, distributed learning is a feasible technique to conquer big data challenges.

This paper aims at error analysis of the distributed learning for (regularization) empirical risk minimization. Given $\mathcal{S} = \{z_i = (\mathbf{x}_i, y_i)\}_{i=1}^N \in (\mathcal{Z} = \mathcal{X} \times \mathcal{Y})^N$, drawn identically and independently from a fixed, but unknown probability distribution \mathbb{P} on $\mathcal{Z} = \mathcal{X} \times \mathcal{Y}$, the (regularization) empirical risk minimization can be stated as

$$\hat{f} = \arg \min_{f \in \mathcal{H}} \hat{R}(f) := \frac{1}{N} \sum_{j=1}^N \ell(f, z_j) + r(f) \quad (1)$$

where $\ell(f, z)$ is a loss function, $r(f)$ is a regularizer, and \mathcal{H} is a hypothesis space. This learning algorithm has been well studied in learning theory, see e.g. [12, 2, 11, 9, 10]. The distributed learning algorithm studied in this paper starts with partitioning the data set \mathcal{S} into m disjoint subsets $\{\mathcal{S}_i\}_{i=1}^m$, $|\mathcal{S}_i| = \frac{N}{m} =: n$. Then it assigns each data subset \mathcal{S}_i to one machine or processor to produce a local estimator \hat{f}_i :

$$\hat{f}_i = \arg \min_{f \in \mathcal{H}} \hat{R}_i(f) := \frac{1}{|\mathcal{S}_i|} \sum_{z_j \in \mathcal{S}_i} \ell(f, z_j) + r(f).$$

The finally global estimator \bar{f} is synthesized by $\bar{f} = \frac{1}{m} \sum_{i=1}^m \hat{f}_i$.

29 Theoretical foundations of distributed learning form a hot topic in machine learning and have been
 30 explored recently in the framework of learning theory [16, 15, 7, 4]. Under local strong convexity,
 31 smoothness and a reasonable set of other conditions, [16] showed that the mean-squared error decays
 32 as

$$\mathbb{E} [\|\bar{f} - f^*\|^2] = \mathcal{O} \left(\frac{1}{N} + \frac{1}{n^2} \right),$$

33 where f^* is the optimal hypothesis in the hypothesis space. Under some eigenfunction assumption,
 34 the error analysis for distributed regularized least squares in reproducing kernel Hilbert space (RKHS)
 35 was established in [15]: if m is not too large,

$$\mathbb{E} [\|\bar{f} - f^*\|^2] = \mathcal{O} \left(\|f_*\|_{\mathcal{H}}^2 + \frac{\gamma(\lambda)}{N} \right),$$

36 where $\gamma(\lambda) = \sum_{j=1}^{\infty} \frac{\mu_j}{\lambda + \mu_j}$, μ_j is the eigenvalue of a Mercer kernel function. Without any eigenfunc-
 37 tion assumption, an improved bound was derived for some $1 \leq p \leq \infty$ [7]:

$$\mathbb{E} [\|\bar{f} - f^*\|_2] = \mathcal{O} \left(\left(\frac{\gamma(\lambda)}{N} \right)^{\frac{1}{2}(1-\frac{1}{p})} \left(\frac{1}{N} \right)^{\frac{1}{2p}} \right).$$

38 There are two main contributions in this paper. First, under strongly convex and smooth, and a
 39 reasonable set of other conditions, we derive a risk bound:

$$R(\bar{f}) - R(f_*) = \mathcal{O} \left(\frac{H_*}{n} + \frac{1}{n^2} - \Delta_{\bar{f}} \right), \quad (2)$$

40 where $R(f) = \mathbb{E}_z[\ell(f, z)] + r(f)$, $\Delta_{\bar{f}} = \mathcal{O} \left(\frac{1}{m^2} \sum_{i,j=1, i \neq j}^m \|\hat{f}_i - \hat{f}_j\|^2 \right)$ is the discrepancy between
 41 all partition-based estimates and $H_* = \mathbb{E}_z[\ell(f_*, z)]$. When the minimal risk is small, i.e., $H_* =$
 42 $\mathcal{O} \left(\frac{1}{n} \right)$, the rate is improved to

$$R(\bar{f}) - R(f_*) = \mathcal{O} \left(\frac{1}{n^2} - \Delta_{\bar{f}} \right).$$

43 Thus, if $m \leq \sqrt{N}$, the order of $R(\bar{f}) - R(f_*)$ is faster than $\mathcal{O} \left(\frac{1}{N} - \Delta_{\bar{f}} \right)$. Note that if $\ell(f, z) + r(f)$
 44 is L -Lipschitz continuous over f , the order of $R(\bar{f}) - R(f^*)$ is

$$R(\bar{f}) - R(f^*) = \mathcal{O} (L \mathbb{E} [\|\bar{f} - f^*\|]) = \mathcal{O} \left(L \sqrt{\mathbb{E} [\|\bar{f} - f^*\|^2]} \right).$$

45 Thus, the order of $R(\bar{f}) - R(f^*)$ in [16, 15, 7] at most $\mathcal{O} \left(\frac{1}{\sqrt{N}} \right)$, which is much slower than that
 46 of our bound $\mathcal{O} \left(\frac{1}{N} \right)$. Our second contribution is to develop a novel max-discrepancy distributed
 47 learning algorithm. From Equation (2), we know that the larger the discrepancy $\Delta_{\bar{f}}$ is, the tighter the
 48 risk bound is. This interesting theoretical finding motivates us to devise a max-discrepancy distributed
 49 learning algorithm (MDD):

$$\hat{f}_i = \arg \min_{f \in \mathcal{H}} \frac{1}{|\mathcal{S}_i|} \sum_{z_j \in \mathcal{S}_i} \ell(f, z_j) + r(f) - \gamma \|f - \bar{f}_{\setminus i}\|_{\mathcal{H}}, \quad (3)$$

50 where $\bar{f}_{\setminus i} = \frac{1}{m-1} \sum_{j=1, j \neq i}^m \hat{f}_j$. The last term of (3) is to make $\Delta_{\bar{f}}$ large. Experimental results on
 51 lots of datasets show that our proposed MDD is sound and efficient.

52 The rest of the paper is organized as follows. In Section 2, we derive a risk bound of distributed
 53 learning with fast convergence rate. In Section 3, we propose two novel algorithms based on the
 54 max-discrepancy of each local estimate in linear space and RKHS. In Section 4, we empirically
 55 analyze the performance of our proposed algorithms. We end in Section 5 with conclusion.

56 2 Error Analysis of Distributed Learning

57 In this section, we will derive a sharper risk bound under some common assumptions.

2.1 Assumptions

In the following, we use $\|\cdot\|_{\mathcal{H}}$ to denote the norm induced by inner product of the Hilbert space \mathcal{H} .
Let the expected risk $R(f)$ and f_* be

$$R(f) = \mathbb{E}_z[\ell(f, z)] + r(f) \text{ and } f_* = \arg \min_{f \in \mathcal{H}} R(f).$$

Assumption 1. The risk $R(f)$ is an η -strongly convex function, that is $\forall f, f' \in \mathcal{H}$,

$$\langle \nabla R(f), f - f' \rangle_{\mathcal{H}} + \frac{\eta}{2} \|f - f'\|_{\mathcal{H}}^2 \leq R(f) - R(f'), \quad (4)$$

or (another equivalent definition) $\forall f, f' \in \mathcal{H}, t \in [0, 1]$,

$$R(tf + (1-t)f') \leq tR(f) + (1-t)R(f') - \frac{1}{2}\eta t(t-1)\|f - f'\|_{\mathcal{H}}^2. \quad (5)$$

Assumption 2. The empirical risk $\hat{R}(f)$ is a convex function.

Assumption 3. The loss function $\ell(f, z)$ is τ -smooth with respect to the first variable f , that is $\forall f, f' \in \mathcal{H}$,

$$\|\nabla \ell(f, \cdot) - \nabla \ell(f', \cdot)\|_{\mathcal{H}} \leq \tau \|f - f'\|_{\mathcal{H}}. \quad (6)$$

Assumption 4. The regularizer $r(f)$ is a τ' -smooth function, that is $\forall f, f' \in \mathcal{H}$,

$$\|\nabla r(f) - \nabla r(f')\|_{\mathcal{H}} \leq \tau' \|f - f'\|_{\mathcal{H}}. \quad (7)$$

Assumption 5. The function $\nu(f, z) = \ell(f, z) + r(f)$ is L -Lipschitz continuous with respect to the first variable f , that is $\forall f, f' \in \mathcal{H}$,

$$\|\nu(f, \cdot) - \nu(f', \cdot)\|_{\mathcal{H}} \leq L \|f - f'\|_{\mathcal{H}}. \quad (8)$$

Assumptions 1, 2, 3, 4 and 5 allow us to model some popular losses, such as square loss and logistic loss, and some regularizer, such as $r(f) = \lambda \|f\|_{\mathcal{H}}^2$.

Assumption 6. We assume that the gradient at f_* is upper bounded by M , that is

$$\|\nabla \ell(f_*, \cdot)\|_{\mathcal{H}} \leq M.$$

Assumption 6 is also a common assumption, which is used in [14, 16].

2.2 Faster Rate of Distributed Learning

Let $\mathcal{N}(\mathcal{H}, \epsilon)$ be the ϵ -net of \mathcal{H} with minimal cardinality, and $C(\mathcal{H}, \epsilon)$ the covering number of $|\mathcal{N}(\mathcal{H}, \epsilon)|$

Theorem 1. For any $0 < \delta < 1$, $\epsilon \geq 0$, under **Assumptions 1, 2, 3, 4, 5 and 6**, and when

$$m \leq \frac{N\eta}{4\tilde{\tau} \log C(\mathcal{H}, \epsilon)}, \quad (9)$$

with probability at least $1 - \delta$, we have

$$\begin{aligned} R(\bar{f}) - R(f_*) &\leq \frac{16\tilde{\tau} \log(4m/\delta)}{n^2\eta} + \frac{128\tau H_* \log(4m/\delta)}{n\eta} + \frac{32\tilde{\tau}^2\epsilon^2}{\eta} + \frac{64\tilde{\tau}L\epsilon \log C(\mathcal{H}, \epsilon)}{n\eta} \\ &\quad + \frac{64\tilde{\tau}\epsilon^2 \log^2 C(\mathcal{H}, \epsilon)}{n^2\eta} - \Delta_{\bar{f}}, \end{aligned} \quad (10)$$

where $\Delta_{\bar{f}} = \frac{\eta}{4m^2} \sum_{i,j=1, i \neq j}^m \|\hat{f}_i - \hat{f}_j\|_{\mathcal{H}}^2$, $H_* = \mathbb{E}_z[\ell(f_*, z)]$ and $\tilde{\tau} = \tau + \tau'$.

From the above theorem, an interesting finding is that, when the larger the discrepancy of each local estimate is, the tighter the risk bound is. Furthermore, one can also see that when ϵ small enough,

$$\frac{32\tilde{\tau}^2\epsilon^2}{\eta} + \frac{64\tilde{\tau}L\epsilon \log C(\mathcal{H}, \epsilon)}{n\eta} + \frac{64\tilde{\tau}\epsilon^2 \log^2 C(\mathcal{H}, \epsilon)}{n^2\eta}$$

will becomes non-dominating. To be specific, we have the following corollary:

80 **Corollary 1.** By setting $\epsilon = \frac{1}{n}$ in Theorem 1, when $m \leq \frac{N\eta}{4\bar{\tau} \log C(\mathcal{H}, 1/n)}$, with high probability, we
 81 have

$$R(\bar{f}) - R(f_*) = \mathcal{O} \left(\frac{H_* \log(m)}{n} + \frac{\log(\mathcal{N}(\mathcal{H}, \frac{1}{n}))}{n^2} - \Delta_{\bar{f}} \right).$$

If the the minimal risk H_* is small, i.e., $H_* = \mathcal{O}(\frac{1}{n})$, the rate can even reach

$$\mathcal{O} \left(\frac{\log(m)}{n^2} + \frac{\log(\mathcal{N}(\mathcal{H}, \frac{1}{n}))}{n^2} - \Delta_{\bar{f}} \right).$$

82 To the best of our knowledge, this is the first $\tilde{\mathcal{O}}(\frac{1}{n^2})$ -type of distributed risk bound for (regularization)
 83 empirical risk minimization.

84 In the next, we will consider two popular Hilbert spaces: linear and reproducing kernel Hilbert space.

85 2.2.1 Linear Space

86 The linear hypothesis space we considered is defined as

$$\mathcal{H} = \left\{ f = \mathbf{w}^T \mathbf{x} \mid \mathbf{w} \in \mathbb{R}^d, \|\mathbf{w}\|_2 \leq B \right\}.$$

87 From [8], the cover number of linear hypothesis space can be bounded by

$$\log(C(\mathcal{H}, \epsilon)) \leq d \log \left(\frac{6B}{\epsilon} \right).$$

88 Thus, if we set $\epsilon = \frac{1}{n}$, from Corollary 1, we have

$$R(\bar{f}) - R(f_*) = \mathcal{O} \left(\frac{H_* \log m}{n} + \frac{d \log n}{n^2} - \Delta_{\bar{f}} \right)$$

89 When the minimal risk is small, i.e., $H_* = \mathcal{O}(\frac{d}{n})$, the rate is improved to

$$\mathcal{O} \left(\frac{d \log(mn)}{n^2} - \Delta_{\bar{f}} \right) = \mathcal{O} \left(\frac{d \log N}{n^2} - \Delta_{\bar{f}} \right).$$

90 Therefore, if $m \leq \sqrt{\frac{N}{d \log N}}$, the order of risk bound can even faster than $\mathcal{O}(\frac{1}{N})$.

91 2.2.2 Reproducing Kernel Hilbert Space

92 The reproducing kernel Hilbert space \mathcal{H}_K associated with the kernel K is defined to be the closure of
 93 the linear span of the set of functions $\{K(\mathbf{x}, \cdot) : \mathbf{x} \in \mathcal{X}\}$ with the inner product satisfying

$$\langle K(\mathbf{x}, \cdot), f \rangle_K = f(\mathbf{x}), \forall \mathbf{x} \in \mathcal{X}, f \in \mathcal{H}_K.$$

94 The hypothesis space of the reproducing kernel Hilbert space we considered in this paper is

$$\mathcal{H} := \{f \in \mathcal{H}_K : \|f\|_K \leq B\}.$$

From [17], if the kernel function K is the popular Gaussian kernel over $[0, 1]^d$:

$$K(\mathbf{x}, \mathbf{x}') = \exp \left\{ -\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{\sigma^2} \right\}, \mathbf{x}, \mathbf{x}' \in [0, 1]^d,$$

then for $0 \leq \epsilon \leq \frac{B}{2}$,

$$\log(C(\mathcal{H}, \epsilon)) = \mathcal{O} \left(\log^d \left(\frac{B}{\epsilon} \right) \right).$$

95 From Corollary 1, if we set $\epsilon = \frac{1}{n}$, and assume $R_* = \mathcal{O}(\frac{1}{n})$, we have

$$R(\bar{f}) - R(f_*) = \mathcal{O} \left(\frac{\log m}{n^2} + \frac{\log^d n}{n^2} - \Delta_{\bar{f}} \right)$$

96 Therefore, if $m \leq \min \left\{ \sqrt{\frac{N}{d \log N}}, \sqrt{\frac{N}{\log^d n}} \right\}$, the order is faster than $\mathcal{O}(\frac{1}{N})$.

97 2.3 Comparison with Related Work

98 In this subsection, we will compare our bound with the related work [16, 15, 7]. Under the smooth,
99 strongly convex and other some assumptions, a distributed risk bound is given in [16]:

$$\mathbb{E} [\|\bar{f} - f_*\|^2] = \mathcal{O} \left(\frac{1}{N} + \frac{\log d}{n^2} \right).$$

100 Under some eigenfunction assumption, the error analysis for distributed regularized least squares
101 were established in [15],

$$\mathbb{E} [\|\bar{f} - f^*\|^2] = \mathcal{O} \left(\|f_*\|_{\mathcal{H}}^2 + \frac{\gamma(\lambda)}{N} \right).$$

102 By removing the eigenfunction assumptions with a novel integral operator method of [15], a new
103 bound was derived [7]:

$$\mathbb{E} [\|\bar{f} - f^*\|] = \mathcal{O} \left(\left(\frac{\gamma(\lambda)}{N} \right)^{\frac{1}{2}(1-\frac{1}{p})} \left(\frac{1}{N} \right)^{\frac{1}{2p}} \right).$$

104 If $\nu(f, z)$ is L -Lipschitz continuous over f , that is

$$\forall f, f' \in \mathcal{H}, z \in \mathcal{Z}, |\nu(f, z) - \nu(f', z)| \leq L\|f - f'\|,$$

105 it is easy to verify that

$$R(f) - R(f_*) \leq L\mathbb{E} [\|\bar{f} - f_*\|] \leq L\sqrt{\mathbb{E} [\|\bar{f} - f_*\|^2]}$$

106 Thus, the order of [16, 15, 7] of $R(f) - R(f_*)$ is at most $\mathcal{O} \left(\frac{1}{\sqrt{N}} \right)$.

107 According to the subsections 2.2.1 and 2.2.2, if m is not very large, and H_* is small, the order of this
108 paper can even faster than $\mathcal{O} \left(\frac{1}{N} \right)$, which is much faster than those of the related work [16, 15, 7].

109 3 Max-Discrepant Distributed Learning (MDD)

110 In this section, we will propose two novel algorithms for linear space and RKHS. From corollary 1,
111 we know that $R(f) - R(f_*) = \mathcal{O} \left(\frac{1}{n^2} - \frac{1}{m^2} \sum_{i,j=1, i \neq j}^m \|\hat{f}_i - \hat{f}_j\|_{\mathcal{H}}^2 \right)$. Thus, to obtain tighter bound,
112 the discrepancy of each local estimate $\hat{f}_i, i = 1, \dots, m$ should be larger.

Algorithm 1 Max-Discrepant Distributed Learning for Linear Space (MDD-LS)

```

1: Input:  $\lambda, \gamma, \mathbf{X}, m, \zeta$ .
2: For each branch node  $i$ :  $\hat{\mathbf{w}}_i^0 = \mathbf{A}_i^{-1} \mathbf{b}_i$ , and push  $\hat{\mathbf{w}}_i^0$  to the center node;
   //  $\mathbf{A}_i = \frac{1}{n} \mathbf{X}_{\mathcal{S}_i} \mathbf{X}_{\mathcal{S}_i}^T + \lambda \mathbf{I}_d$ ,  $\mathbf{b}_i = \frac{1}{n} \mathbf{X}_{\mathcal{S}_i} \mathbf{y}_{\mathcal{S}_i}$ 
3: Center node:  $\bar{\mathbf{w}}^0 = \frac{1}{m} \sum_{i=1}^m \hat{\mathbf{w}}_i^0$  and push  $\bar{\mathbf{w}}_{\setminus i}^0 = \frac{m\bar{\mathbf{w}}^0 - \hat{\mathbf{w}}_i^0}{m-1}$  to each branch node  $i$ ;
4: for  $t = 1, 2, \dots$  do
5:   For each branch node  $i$ :
6:      $\mathbf{d}_i^t = \left( \left( \bar{\mathbf{w}}_{\setminus i}^{t-1} \right)^T \hat{\mathbf{w}}_i^0 \right) / \mathbf{b}_i$ ,  $\hat{\mathbf{w}}_i^t = \hat{\mathbf{w}}_i^0 - \gamma \mathbf{d}_i^t$ ;
7:     push  $\hat{\mathbf{w}}_i^t$  to the center node;
8:   Center node:
9:      $\bar{\mathbf{w}}^t = \frac{1}{m} \sum_{i=1}^m \hat{\mathbf{w}}_i^t$ 
     if  $\|\bar{\mathbf{w}}^t - \bar{\mathbf{w}}^{t-1}\| \leq \zeta$  end for
10:  else
11:    push  $\bar{\mathbf{w}}_{\setminus i}^t = \frac{m\bar{\mathbf{w}}^t - \hat{\mathbf{w}}_i^t}{m-1}$  to each branch node  $i$ 
12:  end for
13: Output:  $\bar{\mathbf{w}} = \frac{1}{m} \sum_{i=1}^m \hat{\mathbf{w}}_i^t$ 

```

113 3.1 Linear Hypothesis Space

114 When \mathcal{H} is a linear Hypothesis space, we consider the following optimization problem:

$$\hat{\mathbf{w}}_i = \arg \min_{\mathbf{w} \in \mathbb{R}^d} \frac{1}{n} \sum_{z_i \in \mathcal{S}_i} (\mathbf{w}^T \mathbf{x}_i - y_i)^2 + \lambda \|\mathbf{w}\|_2^2 + \gamma \mathbf{w}^T \bar{\mathbf{w}}_{\setminus i}, \quad (11)$$

115 where $\bar{\mathbf{w}}_{\setminus i} = \frac{1}{m-1} \sum_{j=1, j \neq i} \hat{\mathbf{w}}_j$. Note that, if given $\bar{\mathbf{w}}_{\setminus i}$, $\hat{\mathbf{w}}_i$ can be written as

$$\hat{\mathbf{w}}_i = \left(\underbrace{\frac{1}{n} \mathbf{X}_{\mathcal{S}_i} \mathbf{X}_{\mathcal{S}_i}^T + \lambda \mathbf{I}_d}_{:= \mathbf{A}_i} \right)^{-1} \left(\underbrace{\frac{1}{n} \mathbf{X}_{\mathcal{S}_i} \mathbf{y}_{\mathcal{S}_i}}_{:= \mathbf{b}_i} - \frac{\gamma \bar{\mathbf{w}}_{\setminus i}}{2} \right),$$

116 where $\mathbf{X}_{\mathcal{S}_i} = (\mathbf{x}_{t_1}, \mathbf{x}_{t_2}, \dots, \mathbf{x}_{t_n})$, $\mathbf{y}_{\mathcal{S}_i} = (y_{t_1}, y_{t_2}, \dots, y_{t_n})^T$, $z_{t_j} \in \mathcal{S}_i$, $j = 1, \dots, n$. In the next,
 117 we will give a iterative algorithm to solve the optimization problem (11), but we should compute
 118 $\mathbf{A}_i^{-1} \bar{\mathbf{w}}_{\setminus i}$, which needs $\mathcal{O}(d^2)$ in each iterative if given \mathbf{A}_i^{-1} , which is computationally intensive.
 119 Fortunately, from Lemma 4 (see in Appendix), the $\mathbf{A}_i^{-1} \bar{\mathbf{w}}_{\setminus i}$ can be computed by

$$\mathbf{A}_i^{-1} \bar{\mathbf{w}}_{\setminus i} = \left(\bar{\mathbf{w}}_{\setminus i}^T \mathbf{c}_i \right) ./ \mathbf{b}_i, \mathbf{c}_i = \mathbf{A}_i^{-1} \mathbf{b}_i$$

120 where $a./c = (a/c_1, \dots, a/c_d)^T$, which only needs $\mathcal{O}(d)$.

121 The Max-Discrepant Distributed Learning algorithm for linear space is given in Algorithm 1. Com-
 122 pared with the traditional divide-and-conquer method, our MDD for linear space only need add $\mathcal{O}(d)$
 in each iteration for each local machine.

Algorithm 2 Max-Discrepant Distributed Learning for RKHS (MDD-RKHS)

- 1: **Input:** λ, γ , kernel function K , \mathbf{X} , m , ζ .
 - 2: For each branch node i : $\hat{\mathbf{w}}_i^0 = \mathbf{A}_i^{-1} \mathbf{b}_i$, and push $\hat{\mathbf{w}}_i^t$ to the center node;
 // $\mathbf{A}_i = \frac{1}{n} \mathbf{K}_{\mathcal{S}_i} + \lambda \mathbf{I}_n$, $\mathbf{b}_i = \frac{1}{n} \mathbf{y}_{\mathcal{S}_i}$
 - 3: Center node: $\hat{\mathbf{g}}_{i,j}^0 = \mathbf{K}_{\mathcal{S}_i, \mathcal{S}_j} \hat{\mathbf{w}}_j^0$, $i, j = 1, \dots, m$, $\bar{\mathbf{g}}_i^0 = \frac{1}{m} \sum_{j=1}^m \hat{\mathbf{g}}_{i,j}^0$ and push $\bar{\mathbf{g}}_i^0 = \frac{m \bar{\mathbf{g}}_i^0 - \hat{\mathbf{g}}_i^0}{m-1}$
 to each branch node i ;
 - 4: **for** $t = 1, 2, \dots$ **do**
 - 5: For each branch node i :
 - 6: $\mathbf{d}_i^t = \left(\left(\bar{\mathbf{g}}_{\setminus i}^{t-1} \right)^T \hat{\mathbf{w}}_i^0 \right) ./ \mathbf{b}_i$, $\hat{\mathbf{w}}_i^t = \hat{\mathbf{w}}_i^0 - \gamma \mathbf{d}_i^t$;
 - 7: push $\hat{\mathbf{w}}_i^t$ to the center node;
 - 8: Center node:
 - 9: $\hat{\mathbf{g}}_{i,j}^t = \mathbf{K}_{\mathcal{S}_i, \mathcal{S}_j} \hat{\mathbf{w}}_j^t$, $i, j = 1, \dots, m$, $\bar{\mathbf{g}}_i^t = \frac{1}{m} \sum_{j=1}^m \hat{\mathbf{g}}_{i,j}^t$
 if $\frac{1}{m} \sum_{i=1}^m \|\bar{\mathbf{g}}_i^t - \bar{\mathbf{g}}_i^{t-1}\| \leq \zeta$ **end for**
 - 10: **else**
 - 11: push $\bar{\mathbf{g}}_{\setminus i}^t = \frac{m \bar{\mathbf{g}}_i^t - \bar{\mathbf{g}}_i^t}{m-1}$ to each branch node i to each branch node i
 - 12: **end for**
 - 13: **Output:** $\bar{f} = \frac{1}{m} \sum_{i=1}^m \hat{f}_i$, where $\hat{f}_i = \mathbf{k}_{\mathcal{S}_i}^T \hat{\mathbf{w}}_i$, where $\mathbf{k}_{\mathcal{S}_i} = (K(\mathbf{x}_1, \cdot), \dots, K(\mathbf{x}_n, \cdot))^T$, $z_j \in \mathcal{S}_i$
-

123

124 3.2 Reproducing Kernel Hilbert Space

125 When \mathcal{H} is a reproducing kernel Hilbert space, that is $f(\mathbf{x}) = \sum_{j=1}^n w_j K(\mathbf{x}_j, \mathbf{x})$, we consider the
 126 following optimization problem:

$$\hat{\mathbf{w}}_i = \arg \min_{\mathbf{w} \in \mathbb{R}^n} \frac{1}{n} \|\mathbf{K}_{\mathcal{S}_i} \mathbf{w} - \mathbf{y}_{\mathcal{S}_i}\|_2^2 + \lambda \mathbf{w}^T \mathbf{K}_{\mathcal{S}_i} \mathbf{w} + \frac{\gamma}{m-1} \sum_{j=1, j \neq i}^m \mathbf{w}^T \mathbf{K}_{\mathcal{S}_i} \mathbf{K}_{\mathcal{S}_i, \mathcal{S}_j} \hat{\mathbf{w}}_j, \quad (12)$$

127 where $\mathbf{K}_{\mathcal{S}_i} = \left[K(\mathbf{x}_{t_j}, \mathbf{x}_{t_{j'}}) \right]_{j, j'=1}^n$, $z_{t_j}, z_{t_{j'}} \in \mathcal{S}_i$, $\mathbf{K}_{\mathcal{S}_i, \mathcal{S}_j} = \left[K(\mathbf{x}_{t_j}, \mathbf{x}_{t_k}) \right]_{j, k=1}^n$, $z_{t_j} \in \mathcal{S}_i, z_{t_k} \in$
 128 \mathcal{S}_j . It is easy to verify that $\hat{\mathbf{w}}_i$ can be written as

$$\hat{\mathbf{w}}_i = \left(\underbrace{\frac{1}{n} \mathbf{K}_{\mathcal{S}_i} + \lambda \mathbf{I}_n}_{:= \mathbf{A}_i} \right)^{-1} \left(\underbrace{\frac{1}{n} \mathbf{y}_{\mathcal{S}_i}}_{:= \mathbf{b}_i} - \frac{\gamma}{2} \bar{\mathbf{g}}_{\setminus i} \right).$$

129 where $\mathbf{g}_j = \mathbf{K}_{S_i, S_j} \hat{\mathbf{w}}_j$ and $\bar{\mathbf{g}}_{\setminus i} = \frac{1}{m-1} \sum_{j=1, j \neq i}^m \hat{\mathbf{g}}_j$.

130 Similar with the linear space, we need to compute $\mathbf{A}_i^{-1} \bar{\mathbf{g}}_{\setminus i}$ in each iterative. From Lemma 4 (see in
131 Appendix), we know that

$$\mathbf{A}_i^{-1} \bar{\mathbf{g}}_{\setminus i} = \left(\bar{\mathbf{g}}_{\setminus i}^T \mathbf{c}_i \right) ./ \mathbf{b}_i, \mathbf{c}_i = \mathbf{A}_i^{-1} \mathbf{b}_i.$$

132 The Max-Discrepant Distributed Learning algorithm for RKHS is also given in Algorithm 2. Com-
133 pared with the traditional divide-and-conquer method, our MDD for RKHS only need add $\mathcal{O}(n)$ in
134 each iteration for local machine.

135 3.3 Complexity

136 **Linear space:** for each node, we need $\mathcal{O}(nd^2)$ to compute the \mathbf{A}_i , $\mathcal{O}(d^3)$ to compute \mathbf{A}_i^{-1} , and
137 $\mathcal{O}(d)$ to compute \mathbf{d}_i for each iterative. Moreover, the communication complexity is $\mathcal{O}(d)$ for each
138 iterative. So, the total complexity is $\mathcal{O}(mnd^2 + md^3 + Tmd)$, where T is the number of iterative.

139 **RKHS:** we need $\mathcal{O}(n^2d)$ to compute the \mathbf{A}_i , $\mathcal{O}(n^3)$ to compute \mathbf{A}_i^{-1} , and $\mathcal{O}(n)$ to compute \mathbf{d}_i for
140 each iterative. Moreover, the communication complexity is $\mathcal{O}(n)$ for each iterative. So, the total
141 complexity is $\mathcal{O}(mn^2d + mn^3 + Tmn)$.

142 **Divide-and-conquer approach:** the complexities of linear space and RKHS are $\mathcal{O}(mnd^2 + md^3)$
143 and $\mathcal{O}(mn^2d + mn^3)$, respectively.

144 4 Experiments

145 In this section, we will compare our MDD methods with the global method and divide-and-conquer
146 method in both Linear Hypothesis and RKHS. Actually, we compare six methods: global Ridge
147 Regression (RR) [5], divide-and-conquer Ridge Regression (DRR) and our MDD-LS (Algorithm 1) in
148 Linear Hypothesis Space, meanwhile, global Kernel Ridge Regression (KRR) [1], divide-and-conquer
149 Kernel Ridge Regression (KDRR) [15] and our MDD-RKHS (Algorithm 2) in Reproducing Kernel
150 Hilbert Space. Based on the recent distributed machine learning platform PARAMETER SERVER
151 [6], we implemented divide-and-conquer methods and MDD methods and do experiments on this
152 platform.

153 We experiment on 10 publicly available dataset from LIBSVM data ¹. We run all methods on a
154 computer node with 32 cores (2.40GHz) and 64 GB memory. While centralized methods only use
155 a single CPU core, distributed methods use all cores to simulate parallel environment. For RKHS
156 methods, we use the popular Gaussian kernels $K(\mathbf{x}, \mathbf{x}') = \exp(-\|\mathbf{x} - \mathbf{x}'\|_2^2 / 2\sigma^2)$ as candidate
157 kernels, and use the best kernel from $\sigma \in \{2^i, i = -10, -9, \dots, 10\}$. The regularized parameterized
158 $\lambda \in \{10^i, i = -6, -5, \dots, 3\}$ in all methods and $\gamma \in \{10^i, i = -6, -5, \dots, 3\}$ in MDD methods
159 are determined by 5-folds cross-validation on training data. For each data set, we run all methods 30
160 times with random partitions on all data sets of non-overlapping 70% training data and 30% testing
161 data, training data.

162 4.1 Performance

163 The root mean square error under fixed 5 worker node (in the next, we will analyze the performance
164 on the different worker nodes) is reported in Table 1, which can be summarized as follows: 1) Global
165 methods outperform the distributed methods on all data sets; 2) In terms of distributed methods,
166 MDD-LS and MDD-RKHS always give better results than the DRR and DKRR. 3) Kernel methods usually
167 give better results than that of Linear methods. 4) Some data sets are sensitive to data partition which
168 result in huge gap between global methods and distributed methods, such as space_ga, phishing and
169 cadata, while others are not.

170 The run time is reported in Table 2, which can be summarized as follows: 1) Global methods cost
171 more time than distributed methods on most data sets. 2) Some centralized methods get closed-form
172 solution in less time, in consideration of file IO and communication cost in distributed methods. 3) In

¹Available at <https://www.csie.ntu.edu.tw/~cjlin/libsvmtools/datasets/>

Table 1: Comparison of average root mean square error of our MDD-LS and MDD-RKHS with RR, DRR, KRR, DKRR.

	madelon	space_ga	cpusmall	phishing	cadata	a8a	a9a	codrna	YearPrediction
c-RR	0.97121	2.58531	45.15010	0.24716	1.93255	0.67161	0.67375	0.84152	12.23623
a-RR	1.40855	2.81430	53.11497	0.25856	2.65991	0.68158	0.68008	0.85555	14.21196
MDD-LS	1.17185	2.67725	46.18480	0.25128	2.11459	0.67730	0.67380	0.84737	12.30163
c-KRR	0.95962	1.45821	53.99303	0.16789	1.50400	0.65961	1.79011	0.67154	/
a-KRR	1.37437	2.38965	54.22893	0.41966	1.59813	0.87312	0.86690	0.67435	5.39791
MDD-RKHS	0.99201	2.03008	54.01543	0.21481	1.55416	0.74578	0.86684	0.67272	5.35078

Table 2: Comparison of run time of our MDD-LS and MDD-RKHS with other methods.

	madelon	space_ga	cpusmall	phishing	cadata	a8a	a9a	codrna	YearPrediction
c-RR	2.06937	0.28045	1.21843	1.52631	0.49022	2.54481	2.95787	1.86626	10.4335
a-RR	1.84997	0.22423	0.46349	0.62537	0.36385	0.77356	0.88100	0.73695	3.70395
MDD-LS	1.87543	0.22424	0.58739	0.66477	0.42771	1.20873	1.16711	0.87666	5.47417
c-KRR	3.45079	1.50828	9.80181	12.0809	76.9912	15.3375	16.1033	137.605	/
a-KRR	2.48753	0.29568	3.37429	1.45185	5.52454	6.02848	5.91322	40.2249	86.7544
MDD-RKHS	2.69209	0.38149	5.63781	1.90103	29.8502	8.62802	9.45444	73.0914	167.208

terms of distributed methods, MDD-LS and MDD-RKHS always need more training time than averaging methods because MDD methods need more than one iteration but averaging methods only need one iteration. 4) Kernel methods always cost more time than Linear methods, because of higher computation complexity. 5) Distributed methods lead huge speedup ratio on some data sets, such as space_ga, phishing and cadata.

The above results show that MDD methods need some additional computation time but make the performance gap between centralized methods and traditional distributed methods tighter, which is consistent with our theoretical analysis.

4.2 Stability

Then, under fixed data set codrna and different worker nodes, we run all methods 30 times with random partitions of non-overlapping 70% training data and 30% testing data, training data were divided into 10 parts every time as well. From Graph [?], we can see 1) The root Mean Square Error grows up very quickly with nodes increasing in the beginning and becomes stable on xxx when more than xxx nodes. 2) The RMSE of MDD methods grows slower than averaging methods. 3) Both of MDD methods and averaging methods convergence on xxx RMSE.

Graph [?] shows 1) The run time decreases fast with nodes increasing at the begin but increasing when nodes size is bigger than xxx. That may caused by training time dominate the run time at beginning but file IO and communication cost play a leading role in following. 2) The run time of MDD methods is always larger than averaging methods and the gap becomes bigger and bigger.

The above results indicates than we can get a trade-off of the root Mean Square Error and run time by controlling worker nodes which can be very useful in practical.

5 Conclusion

In this paper, we studied the generalization performance of distributed learning, and derived a sharper generalization error bound, which is much sharper than existing generalization bounds of divide-and-conquer based distributed learning. Then, we designed two algorithms with statistical guarantees and fast convergence rates for linear space and RKHS: MDD-LS and MDD-RKHS. Empirical results show our methods outperform the popular divide-and-conquer method but only with little additional time.

References

- [1] S. An, W. Liu, and S. Venkatesh. Fast cross-validation algorithms for least squares support vector machine and kernel ridge regression. *Pattern Recognition*, 40(8):2154–2162, 2007.

- 203 [2] A. Caponnetto and E. D. Vito. Optimal rates for the regularized least-squares algorithm.
204 *Foundations of Computational Mathematics*, 7(3):331–368, 2007.
- 205 [3] D. Gillick, A. Faria, and J. DeNero. Mapreduce: Distributed computing for machine learning.
206 *Berkley, Dec*, 18, 2006.
- 207 [4] Z.-C. Guo, S.-B. Lin, and D.-X. Zhou. Learning theory of distributed spectral algorithms.
208 *Inverse Problems*, 33(7):074009, 2017.
- 209 [5] A. E. Hoerl and R. W. Kennard. Ridge regression: Biased estimation for nonorthogonal
210 problems. *Technometrics*, 12(1):55–67, 1970.
- 211 [6] M. Li, D. G. Andersen, J. W. Park, A. J. Smola, A. Ahmed, V. Josifovski, J. Long, E. J.
212 Shekita, and B.-Y. Su. Scaling distributed machine learning with the parameter server. In *OSDI*,
213 volume 14, pages 583–598, 2014.
- 214 [7] S.-B. Lin, X. Guo, and D.-X. Zhou. Distributed learning with regularized least squares. *The*
215 *Journal of Machine Learning Research*, 18(1):3202–3232, 2017.
- 216 [8] G. Pisier. *The volume of convex bodies and Banach space geometry*, volume 94. Cambridge
217 University Press, 1999.
- 218 [9] S. Smale and D.-X. Zhou. Learning theory estimates via integral operators and their approxima-
219 tions. *Constructive approximation*, 26(2):153–172, 2007.
- 220 [10] I. Steinwart and A. Christmann. *Support vector machines*. Springer Verlag, New York, 2008.
- 221 [11] I. Steinwart, D. Hus, and C. Scovel. Optimal rates for regularized least squares regression. In
222 *Proceedings of the Conference on Learning Theory (COLT 2009)*, 2009.
- 223 [12] E. D. Vito, A. Caponnetto, and L. Rosasco. Model selection for regularized least-squares
224 algorithm in learning theory. *Foundations of Computational Mathematics*, 5(1):59–85, 2005.
- 225 [13] X. Wu, X. Zhu, G.-Q. Wu, and W. Ding. Data mining with big data. *IEEE transactions on*
226 *knowledge and data engineering*, 26(1):97–107, 2014.
- 227 [14] L. Zhang, T. Yang, and R. Jin. Empirical risk minimization for stochastic convex optimization:
228 $O(1/n)$ -and $O(1/n^2)$ -type of risk bounds. In *Proceedings of the Conference on Learning*
229 *Theory (COLT 2017)*, pages 1954–1979, 2017.
- 230 [15] Y. Zhang, J. Duchi, and M. Wainwright. Divide and conquer kernel ridge regression. In
231 *Proceedings of Conference on Learning Theory (COLT 2013)*, pages 592–617, 2013.
- 232 [16] Y. Zhang, M. J. Wainwright, and J. C. Duchi. Communication-efficient algorithms for statistical
233 optimization. In *Advances in Neural Information Processing Systems*, pages 1502–1510, 2012.
- 234 [17] D.-X. Zhou. The covering number in learning theory. *Journal of Complexity*, 18(3):739–767,
235 2002.
- 236 [18] Z.-H. Zhou, N. V. Chawla, Y. Jin, and G. J. Williams. Big data opportunities and challenges: Dis-
237 cussions from data analytics perspectives [discussion forum]. *IEEE Computational Intelligence*
238 *Magazine*, 9(4):62–74, 2014.