
Max-Discrepancy Distributed Learning: Fast Risk Bounds and Algorithms

Anonymous Author(s)

Affiliation

Address

email

Abstract

We study the risk performance of distributed learning for the regularization empirical risk minimization with fast convergence rate, substantially improving the error analysis of the existing divide-and-conquer based distributed learning. An interesting theoretical finding is that the larger the discrepancy of each local estimate is, the tighter the risk bound is. This theoretical analysis motivates us to devise an effective max-discrepancy distributed learning algorithm (MDD). Experimental results show that our proposed method can outperform the existing divide-and-conquer methods but with a bit more time. Theoretical analysis and empirical results demonstrate that our MDD is sound and effective.

1 Introduction

In the era of big data, the rapid expansion of computing capacities in automatic data generation and acquisition brings data of unprecedented size and complexity, and raises a series of scientific challenges such as storage bottleneck and algorithmic scalability [18, 15, 7]. Distributed learning based on a divide-and-conquer approach has triggered enormous recent research activities in various areas such as optimization [16] data mining [13] and machine learning [3]. This learning strategy breaks up a big problem into manageable pieces, operates learning algorithms on each piece on individual machines or processors, and then puts the individual solutions together to get a final global output. In this way, distributed learning is a feasible technique to conquer big data challenges.

This paper aims at error analysis of the distributed learning for (regularization) empirical risk minimization. Given $\mathcal{S} = \{z_i = (\mathbf{x}_i, y_i)\}_{i=1}^N \in (\mathcal{Z} = \mathcal{X} \times \mathcal{Y})^N$, drawn identically and independently from a fixed, but unknown probability distribution \mathbb{P} on $\mathcal{Z} = \mathcal{X} \times \mathcal{Y}$, the (regularization) empirical risk minimization can be stated as

$$\hat{f} = \arg \min_{f \in \mathcal{H}} \hat{R}(f) := \frac{1}{N} \sum_{j=1}^N \ell(f, z_j) + r(f) \quad (1)$$

where $\ell(f, z)$ is a loss function, $r(f)$ is a regularizer, and \mathcal{H} is a hypothesis space. This learning algorithm has been well studied in learning theory, see e.g. [12, 2, 11, 9, 10]. The distributed learning algorithm studied in this paper starts with partitioning the data set \mathcal{S} into m disjoint subsets $\{\mathcal{S}_i\}_{i=1}^m$, $|\mathcal{S}_i| = \frac{N}{m} =: n$. Then it assigns each data subset \mathcal{S}_i to one machine or processor to produce a local estimator \hat{f}_i :

$$\hat{f}_i = \arg \min_{f \in \mathcal{H}} \hat{R}_i(f) := \frac{1}{|\mathcal{S}_i|} \sum_{z_j \in \mathcal{S}_i} \ell(f, z_j) + r(f).$$

The finally global estimator \bar{f} is synthesized by $\bar{f} = \frac{1}{m} \sum_{i=1}^m \hat{f}_i$.

29 Theoretical foundations of distributed learning form a hot topic in machine learning and have been
 30 explored recently in the framework of learning theory [16, 15, 7, 4]. Under local strong convexity,
 31 smoothness and a reasonable set of other conditions, [16] showed that the mean-squared error decays
 32 as

$$\mathbb{E} [\|\bar{f} - f^*\|^2] = \mathcal{O} \left(\frac{1}{N} + \frac{1}{n^2} \right),$$

33 where f^* is the optimal hypothesis in the hypothesis space. Under some eigenfunction assumption,
 34 the error analysis for distributed regularized least squares in reproducing kernel Hilbert space (RKHS)
 35 was established in [15]: if m is not too large,

$$\mathbb{E} [\|\bar{f} - f^*\|^2] = \mathcal{O} \left(\|f_*\|_{\mathcal{H}}^2 + \frac{\gamma(\lambda)}{N} \right),$$

36 where $\gamma(\lambda) = \sum_{j=1}^{\infty} \frac{\mu_j}{\lambda + \mu_j}$, μ_j is the eigenvalue of a Mercer kernel function. Without any eigenfunc-
 37 tion assumption, an improved bound was derived for some $1 \leq p \leq \infty$ [7]:

$$\mathbb{E} [\|\bar{f} - f^*\|_2] = \mathcal{O} \left(\left(\frac{\gamma(\lambda)}{N} \right)^{\frac{1}{2}(1-\frac{1}{p})} \left(\frac{1}{N} \right)^{\frac{1}{2p}} \right).$$

38 There are two main contributions in this paper. First, under strongly convex and smooth, and a
 39 reasonable set of other conditions, we derive a risk bound of fast rate:

$$R(\bar{f}) - R(f_*) = \mathcal{O} \left(\frac{H_*}{n} + \frac{1}{n^2} - \Delta_{\bar{f}} \right), \quad (2)$$

40 where $R(f) = \mathbb{E}_z[\ell(f, z)] + r(f)$, $\Delta_{\bar{f}} = \mathcal{O} \left(\frac{1}{m^2} \sum_{i,j=1, i \neq j}^m \|\hat{f}_i - \hat{f}_j\|^2 \right)$ is the discrepancy between
 41 all partition-based estimates and $H_* = \mathbb{E}_z[\ell(f_*, z)]$. When the minimal risk is small, i.e., $H_* =$
 42 $\mathcal{O} \left(\frac{1}{n} \right)$, the rate is improved to

$$R(\bar{f}) - R(f_*) = \mathcal{O} \left(\frac{1}{n^2} - \Delta_{\bar{f}} \right).$$

43 Thus, if $m \leq \sqrt{N}$, the order of $R(\bar{f}) - R(f_*)$ is faster than $\mathcal{O} \left(\frac{1}{N} - \Delta_{\bar{f}} \right)$. Note that if $\ell(f, z) + r(f)$
 44 is L -Lipschitz continuous over f , the order of $R(\bar{f}) - R(f_*)$ is

$$R(\bar{f}) - R(f_*) = \mathcal{O} (L \mathbb{E} [\|\bar{f} - f^*\|]) = \mathcal{O} \left(L \sqrt{\mathbb{E} [\|\bar{f} - f^*\|^2]} \right).$$

45 Thus, the order of $R(\bar{f}) - R(f_*)$ in [16, 15, 7] at most $\mathcal{O} \left(\frac{1}{\sqrt{N}} \right)$, which is much slower than that
 46 of our bound $\mathcal{O} \left(\frac{1}{N} \right)$. Our second contribution is to develop a novel max-discrepancy distributed
 47 learning algorithm. From Equation (2), we know that the larger the discrepancy $\Delta_{\bar{f}}$ is, the tighter the
 48 risk bound is. This interesting theoretical finding motivates us to devise a max-discrepancy distributed
 49 learning algorithm (MDD):

$$\hat{f}_i = \arg \min_{f \in \mathcal{H}} \frac{1}{|\mathcal{S}_i|} \sum_{z_j \in \mathcal{S}_i} \ell(f, z_j) + r(f) - \gamma \|f - \bar{f}_{\setminus i}\|_{\mathcal{H}}, \quad (3)$$

50 where $\bar{f}_{\setminus i} = \frac{1}{m-1} \sum_{j=1, j \neq i}^m \hat{f}_j$. The last term of (3) is to make $\Delta_{\bar{f}}$ large. Experimental results on
 51 lots of datasets show that our proposed MDD is sound and efficient.

52 The rest of the paper is organized as follows. In Section 2, we derive a risk bound of distributed
 53 learning with fast convergence rate. In Section 3, we propose two novel algorithms based on the
 54 max-discrepancy of each local estimate in linear space and RKHS. In Section 4, we empirically
 55 analyze the performance of our proposed algorithms. We end in Section 5 with conclusion. All the
 56 proofs are given in the supplementary material.

57 2 Error Analysis of Distributed Learning

58 In this section, we will derive a sharper risk bound under some common assumptions.

59 2.1 Assumptions

60 In the following, we use $\|\cdot\|_{\mathcal{H}}$ to denote the norm induced by inner product of the Hilbert space \mathcal{H} .
 61 Let the expected risk $R(f)$ and f_* be

$$R(f) = \mathbb{E}_z[\ell(f, z)] + r(f) \text{ and } f_* = \arg \min_{f \in \mathcal{H}} R(f).$$

62 **Assumption 1.** The risk $R(f)$ is an η -strongly convex function, that is $\forall f, f' \in \mathcal{H}$,

$$\langle \nabla R(f), f - f' \rangle_{\mathcal{H}} + \frac{\eta}{2} \|f - f'\|_{\mathcal{H}}^2 \leq R(f) - R(f'), \quad (4)$$

63 or (another equivalent definition) $\forall f, f' \in \mathcal{H}, t \in [0, 1]$,

$$R(tf + (1-t)f') \leq tR(f) + (1-t)R(f') - \frac{1}{2}\eta t(t-1)\|f - f'\|_{\mathcal{H}}^2. \quad (5)$$

64 **Assumption 2.** The empirical risk $\hat{R}(f)$ is a convex function.

65 **Assumption 3.** The loss function $\ell(f, z)$ is τ -smooth with respect to the first variable f , that is
 66 $\forall f, f' \in \mathcal{H}$,

$$\|\nabla \ell(f, \cdot) - \nabla \ell(f', \cdot)\|_{\mathcal{H}} \leq \tau \|f - f'\|_{\mathcal{H}}. \quad (6)$$

67 **Assumption 4.** The regularizer $r(f)$ is a τ' -smooth function, that is $\forall f, f' \in \mathcal{H}$,

$$\|\nabla r(f) - \nabla r(f')\|_{\mathcal{H}} \leq \tau' \|f - f'\|_{\mathcal{H}}. \quad (7)$$

68 **Assumption 5.** The function $\nu(f, z) = \ell(f, z) + r(f)$ is L -Lipschitz continuous with respect to the
 69 first variable f , that is $\forall f, f' \in \mathcal{H}$,

$$\|\nu(f, \cdot) - \nu(f', \cdot)\|_{\mathcal{H}} \leq L \|f - f'\|_{\mathcal{H}}. \quad (8)$$

70 **Assumptions 1, 2, 3, 4 and 5** allow us to model some popular losses, such as square loss and logistic
 71 loss, and some regularizer, such as $r(f) = \lambda \|f\|_{\mathcal{H}}^2$.

72 **Assumption 6.** We assume that the gradient at f_* is upper bounded by M , that is

$$\|\nabla \ell(f_*, \cdot)\|_{\mathcal{H}} \leq M.$$

73 Assumption 6 is also a common assumption, which is used in [14, 16].

74 2.2 Faster Rate of Distributed Learning

75 Let $\mathcal{N}(\mathcal{H}, \epsilon)$ be the ϵ -net of \mathcal{H} with minimal cardinality, and $C(\mathcal{H}, \epsilon)$ the covering number of
 76 $|\mathcal{N}(\mathcal{H}, \epsilon)|$

77 **Theorem 1.** For any $0 < \delta < 1$, $\epsilon \geq 0$, under **Assumptions 1, 2, 3, 4, 5 and 6**, and when

$$m \leq \frac{N\eta}{4\tilde{\tau} \log C(\mathcal{H}, \epsilon)}, \quad (9)$$

78 with probability at least $1 - \delta$, we have

$$\begin{aligned} R(\bar{f}) - R(f_*) &\leq \frac{16\tilde{\tau} \log(4m/\delta)}{n^2\eta} + \frac{128\tau H_* \log(4m/\delta)}{n\eta} + \frac{32\tilde{\tau}^2\epsilon^2}{\eta} + \frac{64\tilde{\tau}L\epsilon \log C(\mathcal{H}, \epsilon)}{n\eta} \\ &\quad + \frac{64\tilde{\tau}\epsilon^2 \log^2 C(\mathcal{H}, \epsilon)}{n^2\eta} - \Delta_{\bar{f}}, \end{aligned} \quad (10)$$

79 where $\Delta_{\bar{f}} = \frac{\eta}{4m^2} \sum_{i,j=1, i \neq j}^m \|\hat{f}_i - \hat{f}_j\|_{\mathcal{H}}^2$, $H_* = \mathbb{E}_z[\ell(f_*, z)]$ and $\tilde{\tau} = \tau + \tau'$.

From the above theorem, an interesting finding is that, when the larger the discrepancy of each local estimate is, the tighter the risk bound is. Furthermore, one can also see that when ϵ small enough,

$$\frac{32\tilde{\tau}^2\epsilon^2}{\eta} + \frac{64\tilde{\tau}L\epsilon \log C(\mathcal{H}, \epsilon)}{n\eta} + \frac{64\tilde{\tau}\epsilon^2 \log^2 C(\mathcal{H}, \epsilon)}{n^2\eta}$$

80 will becomes non-dominating. To be specific, we have the following corollary:

81 **Corollary 1.** By setting $\epsilon = \frac{1}{n}$ in Theorem 1, when $m \leq \frac{N\eta}{4\bar{\tau} \log C(\mathcal{H}, 1/n)}$, with high probability, we
 82 have

$$R(\bar{f}) - R(f_*) = \mathcal{O} \left(\frac{H_* \log(m)}{n} + \frac{\log(\mathcal{N}(\mathcal{H}, \frac{1}{n}))}{n^2} - \Delta_{\bar{f}} \right).$$

If the the minimal risk H_* is small, i.e., $H_* = \mathcal{O}(\frac{1}{n})$, the rate can reach

$$\mathcal{O} \left(\frac{\log(m)}{n^2} + \frac{\log(\mathcal{N}(\mathcal{H}, \frac{1}{n}))}{n^2} - \Delta_{\bar{f}} \right).$$

83 To the best of our knowledge, this is the first $\tilde{\mathcal{O}}(\frac{1}{n^2})$ -type of distributed risk bound for (regularization)
 84 empirical risk minimization.

85 In the next, we will consider two popular Hilbert spaces: linear and reproducing kernel Hilbert space.

86 2.2.1 Linear Space

87 The linear hypothesis space we considered is defined as

$$\mathcal{H} = \left\{ f = \mathbf{w}^T \mathbf{x} \mid \mathbf{w} \in \mathbb{R}^d, \|\mathbf{w}\|_2 \leq B \right\}.$$

88 From [8], the cover number of linear hypothesis space can be bounded by

$$\log(C(\mathcal{H}, \epsilon)) \leq d \log \left(\frac{6B}{\epsilon} \right).$$

89 Thus, if we set $\epsilon = \frac{1}{n}$, from Corollary 1, we have

$$R(\bar{f}) - R(f_*) = \mathcal{O} \left(\frac{H_* \log m}{n} + \frac{d \log n}{n^2} - \Delta_{\bar{f}} \right)$$

90 When the minimal risk is small, i.e., $H_* = \mathcal{O}(\frac{d}{n})$, the rate is improved to

$$\mathcal{O} \left(\frac{d \log(mn)}{n^2} - \Delta_{\bar{f}} \right) = \mathcal{O} \left(\frac{d \log N}{n^2} - \Delta_{\bar{f}} \right).$$

91 Therefore, if $m \leq \sqrt{\frac{N}{d \log N}}$, the order of risk bound can even faster than $\mathcal{O}(\frac{1}{N})$.

92 2.2.2 Reproducing Kernel Hilbert Space

93 The reproducing kernel Hilbert space \mathcal{H}_K associated with the kernel K is defined to be the closure of
 94 the linear span of the set of functions $\{K(\mathbf{x}, \cdot) : \mathbf{x} \in \mathcal{X}\}$ with the inner product satisfying

$$\langle K(\mathbf{x}, \cdot), f \rangle_K = f(\mathbf{x}), \forall \mathbf{x} \in \mathcal{X}, f \in \mathcal{H}_K.$$

95 The hypothesis space of the reproducing kernel Hilbert space we considered in this paper is

$$\mathcal{H} := \{f \in \mathcal{H}_K : \|f\|_K \leq B\}.$$

From [17], if the kernel function K is the popular Gaussian kernel over $[0, 1]^d$:

$$K(\mathbf{x}, \mathbf{x}') = \exp \left\{ -\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{\sigma^2} \right\}, \mathbf{x}, \mathbf{x}' \in [0, 1]^d,$$

then for $0 \leq \epsilon \leq \frac{B}{2}$,

$$\log(C(\mathcal{H}, \epsilon)) = \mathcal{O} \left(\log^d \left(\frac{B}{\epsilon} \right) \right).$$

96 From Corollary 1, if we set $\epsilon = \frac{1}{n}$, and assume $R_* = \mathcal{O}(\frac{1}{n})$, we have

$$R(\bar{f}) - R(f_*) = \mathcal{O} \left(\frac{\log m}{n^2} + \frac{\log^d n}{n^2} - \Delta_{\bar{f}} \right)$$

97 Therefore, if $m \leq \min \left\{ \sqrt{\frac{N}{d \log N}}, \sqrt{\frac{N}{\log^d n}} \right\}$, the order is faster than $\mathcal{O}(\frac{1}{N})$.

98 2.3 Comparison with Related Work

99 In this subsection, we will compare our bound with the related work [16, 15, 7]. Under the smooth,
100 strongly convex and other some assumptions, a distributed risk bound is given in [16]:

$$\mathbb{E} [\|\bar{f} - f_*\|^2] = \mathcal{O} \left(\frac{1}{N} + \frac{\log d}{n^2} \right).$$

101 Under some eigenfunction assumption, the error analysis for distributed regularized least squares
102 were established in [15],

$$\mathbb{E} [\|\bar{f} - f_*\|^2] = \mathcal{O} \left(\|f_*\|_{\mathcal{H}}^2 + \frac{\gamma(\lambda)}{N} \right).$$

103 By removing the eigenfunction assumptions with a novel integral operator method of [15], a new
104 bound was derived [7]:

$$\mathbb{E} [\|\bar{f} - f_*\|] = \mathcal{O} \left(\left(\frac{\gamma(\lambda)}{N} \right)^{\frac{1}{2}(1-\frac{1}{p})} \left(\frac{1}{N} \right)^{\frac{1}{2p}} \right).$$

105 If $\nu(f, z)$ is L -Lipschitz continuous over f , that is

$$\forall f, f' \in \mathcal{H}, z \in \mathcal{Z}, |\nu(f, z) - \nu(f', z)| \leq L\|f - f'\|,$$

106 it is easy to verify that

$$R(f) - R(f_*) \leq L\mathbb{E} [\|\bar{f} - f_*\|] \leq L\sqrt{\mathbb{E} [\|\bar{f} - f_*\|^2]}$$

107 Thus, the order of [16, 15, 7] of $R(f) - R(f_*)$ is at most $\mathcal{O} \left(\frac{1}{\sqrt{N}} \right)$.

108 According to the subsections 2.2.1 and 2.2.2, if m is not very large, and H_* is small, the order of this
109 paper can even faster than $\mathcal{O} \left(\frac{1}{N} \right)$, which is much faster than those of the related work [16, 15, 7].

110 3 Max-Discrepant Distributed Learning (MDD)

111 In this section, we will propose two novel algorithms for linear space and RKHS. From corollary 1,
112 we know that $R(f) - R(f_*) = \mathcal{O} \left(\frac{1}{n^2} - \frac{1}{m^2} \sum_{i,j=1, i \neq j}^m \|\hat{f}_i - \hat{f}_j\|_{\mathcal{H}}^2 \right)$. Thus, to obtain tighter bound,
113 the discrepancy of each local estimate $\hat{f}_i, i = 1, \dots, m$ should be larger.

Algorithm 1 Max-Discrepant Distributed Learning for Linear Space (MDD-LS)

```

1: Input:  $\lambda, \gamma, \mathbf{X}, m, \zeta$ .
2: For each worker node  $i$ :  $\hat{\mathbf{w}}_i^0 = \mathbf{A}_i^{-1} \mathbf{b}_i$ , and push  $\hat{\mathbf{w}}_i^0$  to the server node.
   //  $\mathbf{A}_i = \frac{1}{n} \mathbf{X}_{\mathcal{S}_i} \mathbf{X}_{\mathcal{S}_i}^T + \lambda \mathbf{I}_d$ ,  $\mathbf{b}_i = \frac{1}{n} \mathbf{X}_{\mathcal{S}_i} \mathbf{y}_{\mathcal{S}_i}$ .
3: For server node:  $\bar{\mathbf{w}}^0 = \frac{1}{m} \sum_{i=1}^m \hat{\mathbf{w}}_i^0$ ,  $\bar{\mathbf{w}}_{\setminus i}^0 = \frac{m\bar{\mathbf{w}}^0 - \hat{\mathbf{w}}_i^0}{m-1}$ .
4: for  $t = 1, 2, \dots$  do
5:   For each worker node  $i$ :
     Pull  $\bar{\mathbf{w}}_{\setminus i}^{t-1}$  from server node.
6:    $\mathbf{d}_i^t = \left( \left( \bar{\mathbf{w}}_{\setminus i}^{t-1} \right)^T \hat{\mathbf{w}}_i^0 \right) / \mathbf{b}_i$ ,  $\hat{\mathbf{w}}_i^t = \hat{\mathbf{w}}_i^0 - \gamma \mathbf{d}_i^t$ .
7:   Push  $\hat{\mathbf{w}}_i^t$  to the server node.
8:   For server node:
9:      $\bar{\mathbf{w}}^t = \frac{1}{m} \sum_{i=1}^m \hat{\mathbf{w}}_i^t$ 
     if  $\|\bar{\mathbf{w}}^t - \bar{\mathbf{w}}^{t-1}\| \leq \zeta$  end for
10:    else  $\bar{\mathbf{w}}_{\setminus i}^t = \frac{m\bar{\mathbf{w}}^t - \hat{\mathbf{w}}_i^t}{m-1}$ .
11:  end for
12: Output:  $\bar{\mathbf{w}} = \frac{1}{m} \sum_{i=1}^m \hat{\mathbf{w}}_i^t$ .

```

114 3.1 Linear Hypothesis Space

115 When \mathcal{H} is a linear Hypothesis space, we consider the following optimization problem:

$$\hat{\mathbf{w}}_i = \arg \min_{\mathbf{w} \in \mathbb{R}^d} \frac{1}{n} \sum_{z_i \in \mathcal{S}_i} (\mathbf{w}^T \mathbf{x}_i - y_i)^2 + \lambda \|\mathbf{w}\|_2^2 + \gamma \mathbf{w}^T \bar{\mathbf{w}}_{\setminus i}, \quad (11)$$

116 where $\bar{\mathbf{w}}_{\setminus i} = \frac{1}{m-1} \sum_{j=1, j \neq i} \hat{\mathbf{w}}_j$. Note that, if given $\bar{\mathbf{w}}_{\setminus i}$, $\hat{\mathbf{w}}_i$ has following closed form solution:

$$\hat{\mathbf{w}}_i = \left(\underbrace{\frac{1}{n} \mathbf{X}_{\mathcal{S}_i} \mathbf{X}_{\mathcal{S}_i}^T + \lambda \mathbf{I}_d}_{:= \mathbf{A}_i} \right)^{-1} \left(\underbrace{\frac{1}{n} \mathbf{X}_{\mathcal{S}_i} \mathbf{y}_{\mathcal{S}_i}}_{:= \mathbf{b}_i} - \frac{\gamma \bar{\mathbf{w}}_{\setminus i}}{2} \right),$$

117 where $\mathbf{X}_{\mathcal{S}_i} = (\mathbf{x}_{t_1}, \mathbf{x}_{t_2}, \dots, \mathbf{x}_{t_n})$, $\mathbf{y}_{\mathcal{S}_i} = (y_{t_1}, y_{t_2}, \dots, y_{t_n})^T$, $z_{t_j} \in \mathcal{S}_i$, $j = 1, \dots, n$. In the next,
 118 we will give a iterative algorithm to solve the optimization problem (11). In each iteration, we should
 119 compute $\mathbf{A}_i^{-1} \bar{\mathbf{w}}_{\setminus i}$, which needs $\mathcal{O}(d^2)$ if given \mathbf{A}_i^{-1} , which is computational intensive. Fortunately,
 120 from Lemma 4 (see in Appendix), the $\mathbf{A}_i^{-1} \bar{\mathbf{w}}_{\setminus i}$ can be computed by

$$\mathbf{A}_i^{-1} \bar{\mathbf{w}}_{\setminus i} = \left(\bar{\mathbf{w}}_{\setminus i}^T \mathbf{c}_i \right) ./ \mathbf{b}_i, \mathbf{c}_i = \mathbf{A}_i^{-1} \mathbf{b}_i$$

121 where $a./c = (a/c_1, \dots, a/c_d)^T$, which only needs $\mathcal{O}(d)$.

122 The Max-Discrepant Distributed Learning algorithm for linear space is given in Algorithm 1. Com-
 123 pared with the traditional divide-and-conquer method, our MDD for linear space only need add $\mathcal{O}(d)$
 in each iteration for each worker node.

Algorithm 2 Max-Discrepant Distributed Learning for RKHS (MDD-RKHS)

```

1: Input:  $\lambda, \gamma$ , kernel function  $K$ ,  $\mathbf{X}$ ,  $m$ ,  $\zeta$ .
2: For each worker node  $i$ :  $\hat{\mathbf{w}}_i^0 = \mathbf{A}_i^{-1} \mathbf{b}_i$ , and push  $\hat{\mathbf{w}}_i^t$  to the server node.
   //  $\mathbf{A}_i = \frac{1}{n} \mathbf{K}_{\mathcal{S}_i} + \lambda \mathbf{I}_n$ ,  $\mathbf{b}_i = \frac{1}{n} \mathbf{y}_{\mathcal{S}_i}$ .
3: For server node:  $\hat{\mathbf{g}}_{i,j}^0 = \mathbf{K}_{\mathcal{S}_i, \mathcal{S}_j} \hat{\mathbf{w}}_j^0$ ,  $i, j = 1, \dots, m$ ,  $\bar{\mathbf{g}}_{\setminus i}^0 = \frac{m \hat{\mathbf{g}}_i^0 - \hat{\mathbf{g}}_i^0}{m-1}$ .
4: for  $t = 1, 2, \dots$  do
5:   For each worker node  $i$ :
6:     Pull  $\bar{\mathbf{g}}_{\setminus i}^{t-1}$  from server node.
7:      $\mathbf{d}_i^t = \left( \left( \bar{\mathbf{g}}_{\setminus i}^{t-1} \right)^T \hat{\mathbf{w}}_i^0 \right) ./ \mathbf{b}_i$ ,  $\hat{\mathbf{w}}_i^t = \hat{\mathbf{w}}_i^0 - \gamma \mathbf{d}_i^t$ .
8:     Push  $\hat{\mathbf{w}}_i^t$  to the server node.
9:   For server node:
10:     $\hat{\mathbf{g}}_{i,j}^t = \mathbf{K}_{\mathcal{S}_i, \mathcal{S}_j} \hat{\mathbf{w}}_j^t$ ,  $i, j = 1, \dots, m$ ,  $\bar{\mathbf{g}}_i^t = \frac{1}{m} \sum_{j=1}^m \hat{\mathbf{g}}_{i,j}^t$ .
    if  $\frac{1}{m} \sum_{i=1}^m \|\bar{\mathbf{g}}_i^t - \bar{\mathbf{g}}_i^t\| \leq \zeta$  end for
11:   else
12:      $\bar{\mathbf{g}}_{\setminus i}^t = \frac{m \bar{\mathbf{g}}_i^t - \bar{\mathbf{g}}_i^t}{m-1}$ .
13:   end for
14: Output:  $\bar{f} = \frac{1}{m} \sum_{i=1}^m \hat{f}_i$ , where  $\hat{f}_i = \mathbf{k}_{\mathcal{S}_i}^T \hat{\mathbf{w}}_i$ , where  $\mathbf{k}_{\mathcal{S}_i} = (K(\mathbf{x}_1, \cdot), \dots, K(\mathbf{x}_n, \cdot))^T$ ,  $z_j \in \mathcal{S}_i$ 
```

124

125 3.2 Reproducing Kernel Hilbert Space

126 When \mathcal{H} is a reproducing kernel Hilbert space, that is $f(\mathbf{x}) = \sum_{j=1}^n w_j K(\mathbf{x}_j, \mathbf{x})$, we consider the
 127 following optimization problem:

$$\hat{\mathbf{w}}_i = \arg \min_{\mathbf{w} \in \mathbb{R}^n} \frac{1}{n} \|\mathbf{K}_{\mathcal{S}_i} \mathbf{w} - \mathbf{y}_{\mathcal{S}_i}\|_2^2 + \lambda \mathbf{w}^T \mathbf{K}_{\mathcal{S}_i} \mathbf{w} + \frac{\gamma}{m-1} \sum_{j=1, j \neq i}^m \mathbf{w}^T \mathbf{K}_{\mathcal{S}_i} \mathbf{K}_{\mathcal{S}_i, \mathcal{S}_j} \hat{\mathbf{w}}_j, \quad (12)$$

where $\mathbf{K}_{S_i} = \left[K(\mathbf{x}_{t_j}, \mathbf{x}_{t_{j'}}) \right]_{j,j'=1}^n$, $z_{t_j}, z_{t_{j'}} \in S_i$, $\mathbf{K}_{S_i, S_j} = \left[K(\mathbf{x}_{t_j}, \mathbf{x}_{t_k}) \right]_{j,k=1}^n$, $z_{t_j} \in S_i, z_{t_k} \in S_j$. It is easy to verify that $\hat{\mathbf{w}}_i$ can be written as

$$\hat{\mathbf{w}}_i = \underbrace{\left(\frac{1}{n} \mathbf{K}_{S_i} + \lambda \mathbf{I}_n \right)}_{:= \mathbf{A}_i}^{-1} \underbrace{\left(\frac{1}{n} \mathbf{y}_{S_i} - \frac{\gamma}{2} \bar{\mathbf{g}}_{\setminus i} \right)}_{:= \mathbf{b}_i}.$$

where $\mathbf{g}_j = \mathbf{K}_{S_i, S_j} \hat{\mathbf{w}}_j$ and $\bar{\mathbf{g}}_{\setminus i} = \frac{1}{m-1} \sum_{j=1, j \neq i}^m \mathbf{g}_j$.

Similar with the linear space, we need to compute $\mathbf{A}_i^{-1} \bar{\mathbf{g}}_{\setminus i}$ in each iterative. From Lemma 4 (see in Appendix), we know that

$$\mathbf{A}_i^{-1} \bar{\mathbf{g}}_{\setminus i} = \left(\bar{\mathbf{g}}_{\setminus i}^T \mathbf{c}_i \right) ./ \mathbf{b}_i, \mathbf{c}_i = \mathbf{A}_i^{-1} \mathbf{b}_i.$$

The Max-Discrepant Distributed Learning algorithm for RKHS is also given in Algorithm 2. Compared with the traditional divide-and-conquer method, our MDD for RKHS only need add $\mathcal{O}(n)$ in each iteration for local machine.

3.3 Complexity

Linear space: At the very beginning, we need $\mathcal{O}(nd^2)$ to compute the \mathbf{A}_i , $\mathcal{O}(d^3)$ to compute \mathbf{A}_i^{-1} for each worker node. In each iteration, worker nodes cost $\mathcal{O}(d)$ to compute \mathbf{d}_i^t and the server node costs $\mathcal{O}(md)$ to compute $\bar{\mathbf{w}}_{\setminus i}^t$. So, the sequential computation complexity is $\mathcal{O}(nd^2 + d^3 + Tmd)$, where T is the number of iteration. Moreover, the total communication complexity is $\mathcal{O}(Td)$.

RKHS: At the very beginning, we need $\mathcal{O}(n^2d)$ to compute the \mathbf{A}_i and $\mathcal{O}(n^3)$ to compute \mathbf{A}_i^{-1} . In each iteration, worker nodes cost $\mathcal{O}(n)$ to compute \mathbf{d}_i^t and the server node costs $\mathcal{O}(mn)$ to compute $\bar{\mathbf{g}}_{\setminus i}^t$. So, the sequential computation complexity is $\mathcal{O}(n^2d + n^3 + Tmn)$, where T is the number of iteration. Moreover, the total communication complexity is $\mathcal{O}(Tn)$.

Divide-and-conquer approach: The sequential complexities of linear space and RKHS are $\mathcal{O}(nd^2 + d^3)$ and $\mathcal{O}(n^2d + n^3)$, respectively. Meanwhile, the communication complexities are $\mathcal{O}(d)$ and $\mathcal{O}(n)$.

4 Experiments

In this section, we will compare our MDD methods with the global method and divide-and-conquer method in both Linear and RKHS Hypothesis. Actually, we compare six approaches: global Ridge Regression (RR) [5], divide-and-conquer Ridge Regression (DRR) and our MDD-LS (Algorithm 1) in Linear Hypothesis Space, meanwhile, global Kernel Ridge Regression (KRR) [1], divide-and-conquer Kernel Ridge Regression (KDRR) [15] and our MDD-RKHS (Algorithm 2) in Reproducing Kernel Hilbert Space. Based on the recent distributed machine learning platform PARAMETER SERVER [6], we implemented divide-and-conquer methods and MDD methods and do experiments on this framework.

We experiment on 10 publicly available datasets from LIBSVM data¹. We run all methods on a computer node with 32 cores (2.40GHz) and 64 GB memory. While global methods only use a single CPU core, distributed methods use all cores to simulate parallel environment. For RKHS methods, we use the popular Gaussian kernels $K(\mathbf{x}, \mathbf{x}') = \exp(-\|\mathbf{x} - \mathbf{x}'\|_2^2 / 2\sigma^2)$ as candidate kernels, and choose the best kernel from $\sigma \in \{2^i, i = -10, -9, \dots, 10\}$ by 5-folds cross-validation. The regularized parameterized $\lambda \in \{10^i, i = -6, -5, \dots, 3\}$ in all methods and $\gamma \in \{10^i, i = -6, -5, \dots, 3\}$ in MDD methods are determined by 5-folds cross-validation on training data. With the same kernel and parameters, for each data set, we run all methods 30 times with random partitions on all data sets of non-overlapping 70% training data and 30% testing data. All statement of statistical significance in the remainder refer to a 95% level of significance under t -test.

The root mean square error of all methods is reported in Table 1. Meanwhile, we repeat distributed methods on different amount of worker nodes, 5 and 10 for simplification. The table can be

¹Available at <https://www.csie.ntu.edu.tw/~cjlin/libsvmtools/datasets/>

Table 1: Comparison of average root mean square error of our MDD-LS and MDD-RKHS with RR, DRR, KRR, DKRR. We bold the numbers of the best method, and underline the numbers of the other methods which are not significantly worse than the best one

	madelon	space_ga	cpusmall	phishing	cadata	a8a	a9a	codrna	YearPred
RR	0.971	2.585	45.150	0.247	1.932	0.671	0.673	0.841	12.233
DRR-5	<u>0.989</u>	2.814	53.114	0.262	2.659	<u>0.681</u>	<u>0.680</u>	0.855	14.216
DRR-10	1.408	2.983	55.557	0.273	2.839	<u>0.725</u>	<u>0.696</u>	0.863	15.780
MDD-LS-5	<u>0.977</u>	<u>2.677</u>	46.184	<u>0.257</u>	<u>2.114</u>	<u>0.677</u>	<u>0.673</u>	<u>0.847</u>	<u>12.303</u>
MDD-LS-10	1.121	2.750	48.956	0.268	<u>2.352</u>	<u>0.703</u>	<u>0.685</u>	<u>0.854</u>	14.158
KRR	0.959	1.458	53.993	0.167	1.504	0.659	0.790	0.671	/
KDRR-5	1.142	2.389	54.228	0.419	1.598	0.873	0.866	0.674	5.397
KDRR-10	1.374	2.531	56.233	0.422	1.824	0.906	0.893	0.687	5.631
MDD-RKHS-5	<u>0.992</u>	2.030	<u>54.015</u>	0.214	<u>1.554</u>	<u>0.745</u>	<u>0.804</u>	<u>0.672</u>	5.350
MDD-RKHS-10	1.292	2.326	55.120	0.239	1.780	<u>0.773</u>	0.849	0.683	5.534

Table 2: Comparison of run time (second) around our proposed MDD-LS and MDD-RKHS with other methods.

	madelon	space_ga	cpusmall	phishing	cadata	a8a	a9a	codrna	YearPred
RR	2.069	0.280	1.218	1.526	0.490	2.544	2.957	1.866	10.433
DRR-5	0.849	0.094	0.463	0.625	0.363	0.773	0.881	0.736	3.709
DRR-10	0.623	0.193	0.298	0.350	0.214	0.401	0.503	0.435	2.645
MDD-LS-5	0.875	0.115	0.587	0.664	0.427	1.208	1.167	0.876	5.474
MDD-LS-10	0.656	0.134	0.315	0.395	0.269	0.651	0.628	0.412	3.156
KRR	3.450	1.508	9.801	12.08	76.99	15.33	16.103	137.6	/
KDRR-5	1.487	0.295	3.374	1.451	5.524	6.021	5.913	40.22	86.754
KDRR-10	0.983	0.183	1.863	0.689	0.302	3.670	3.544	23.64	46.197
MDD-RKHS-5	1.692	0.331	5.637	1.901	29.85	8.628	9.454	73.09	167.208
MDD-RKHS-10	1.041	0.206	3.024	0.984	17.78	4.125	5.679	40.23	89.312

summarized as follows: 1) Our MDD-LS and MDD-RKHS exhibits better prediction accuracy than the DRR and KDRR over almost all data sets for both linear space and RKHS. This demonstrates the advantage of MDD methods in generalization performance. 2) Our MDD-LS and MDD-RKHS give comparable result with global methods on most of data sets. 3) Kernel methods can usually get more optimal results than that of Linear methods; 4) Some data sets are sensitive to data partition, whose results existing huge gap between global methods and distributed methods, such as space_ga and phishing for RKHS, while others are not; 5) The increase of worker nodes causes higher root mean square error.

The run time is reported in Table 2, which can be summarized as follows: 1) Global methods cost more time than distributed methods on all data sets; 2) Kernel methods always spend more time than Linear methods, because of higher computation complexity; 3) Distributed methods lead great speedup on some data sets, such as space_ga, phishing and cadata; 4) The running time of distributed methods decays almost linearly associated with the increase of worker nodes; 5) Compared with global methods, our MDD methods own higher computational efficiency, while existing small distance away from divide-and-conquer methods.

The above results show that MDD methods need a bit more training time but make the performance gap between global methods and traditional distributed methods tighter, which is consistent with our theoretical analysis.

Note: Source code are attached in Supplementary Material.

5 Conclusion

In this paper, we studied the generalization performance of distributed learning, and derived a sharper generalization error bound, which is much sharper than existing generalization bounds of divide-and-conquer based distributed learning. Then, we designed two algorithms with statistical guarantees and fast convergence rates for linear space and RKHS: MDD-LS and MDD-RKHS. Empirical results show our methods outperform the popular divide-and-conquer method but only with a bit more time.

References

- [1] S. An, W. Liu, and S. Venkatesh. Fast cross-validation algorithms for least squares support vector machine and kernel ridge regression. *Pattern Recognition*, 40(8):2154–2162, 2007.
- [2] A. Caponnetto and E. D. Vito. Optimal rates for the regularized least-squares algorithm. *Foundations of Computational Mathematics*, 7(3):331–368, 2007.
- [3] D. Gillick, A. Faria, and J. DeNero. Mapreduce: Distributed computing for machine learning. *Berkley, Dec*, 18, 2006.
- [4] Z.-C. Guo, S.-B. Lin, and D.-X. Zhou. Learning theory of distributed spectral algorithms. *Inverse Problems*, 33(7):074009, 2017.
- [5] A. E. Hoerl and R. W. Kennard. Ridge regression: Biased estimation for nonorthogonal problems. *Technometrics*, 12(1):55–67, 1970.
- [6] M. Li, D. G. Andersen, J. W. Park, A. J. Smola, A. Ahmed, V. Josifovski, J. Long, E. J. Shekita, and B.-Y. Su. Scaling distributed machine learning with the parameter server. In *OSDI*, volume 14, pages 583–598, 2014.
- [7] S.-B. Lin, X. Guo, and D.-X. Zhou. Distributed learning with regularized least squares. *The Journal of Machine Learning Research*, 18(1):3202–3232, 2017.
- [8] G. Pisier. *The volume of convex bodies and Banach space geometry*, volume 94. Cambridge University Press, 1999.
- [9] S. Smale and D.-X. Zhou. Learning theory estimates via integral operators and their approximations. *Constructive approximation*, 26(2):153–172, 2007.
- [10] I. Steinwart and A. Christmann. *Support vector machines*. Springer Verlag, New York, 2008.
- [11] I. Steinwart, D. Hus, and C. Scovel. Optimal rates for regularized least squares regression. In *Proceedings of the Conference on Learning Theory (COLT 2009)*, 2009.
- [12] E. D. Vito, A. Caponnetto, and L. Rosasco. Model selection for regularized least-squares algorithm in learning theory. *Foundations of Computational Mathematics*, 5(1):59–85, 2005.
- [13] X. Wu, X. Zhu, G.-Q. Wu, and W. Ding. Data mining with big data. *IEEE transactions on knowledge and data engineering*, 26(1):97–107, 2014.
- [14] L. Zhang, T. Yang, and R. Jin. Empirical risk minimization for stochastic convex optimization: $O(1/n)$ -and $O(1/n^2)$ -type of risk bounds. In *Proceedings of the Conference on Learning Theory (COLT 2017)*, pages 1954–1979, 2017.
- [15] Y. Zhang, J. Duchi, and M. Wainwright. Divide and conquer kernel ridge regression. In *Proceedings of Conference on Learning Theory (COLT 2013)*, pages 592–617, 2013.
- [16] Y. Zhang, M. J. Wainwright, and J. C. Duchi. Communication-efficient algorithms for statistical optimization. In *Advances in Neural Information Processing Systems*, pages 1502–1510, 2012.
- [17] D.-X. Zhou. The covering number in learning theory. *Journal of Complexity*, 18(3):739–767, 2002.
- [18] Z.-H. Zhou, N. V. Chawla, Y. Jin, and G. J. Williams. Big data opportunities and challenges: Discussions from data analytics perspectives [discussion forum]. *IEEE Computational Intelligence Magazine*, 9(4):62–74, 2014.